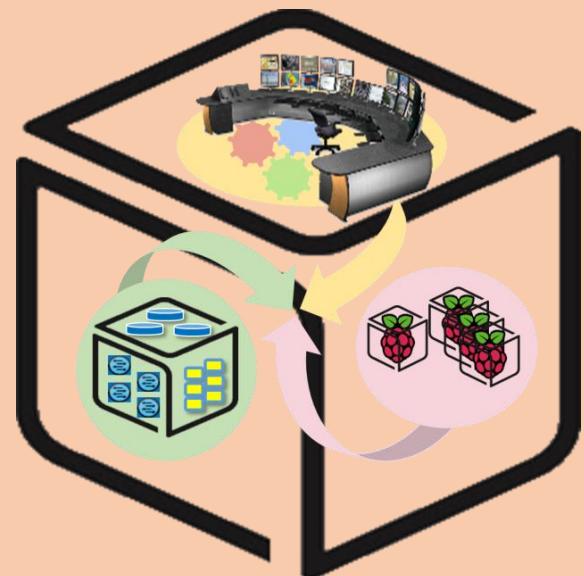


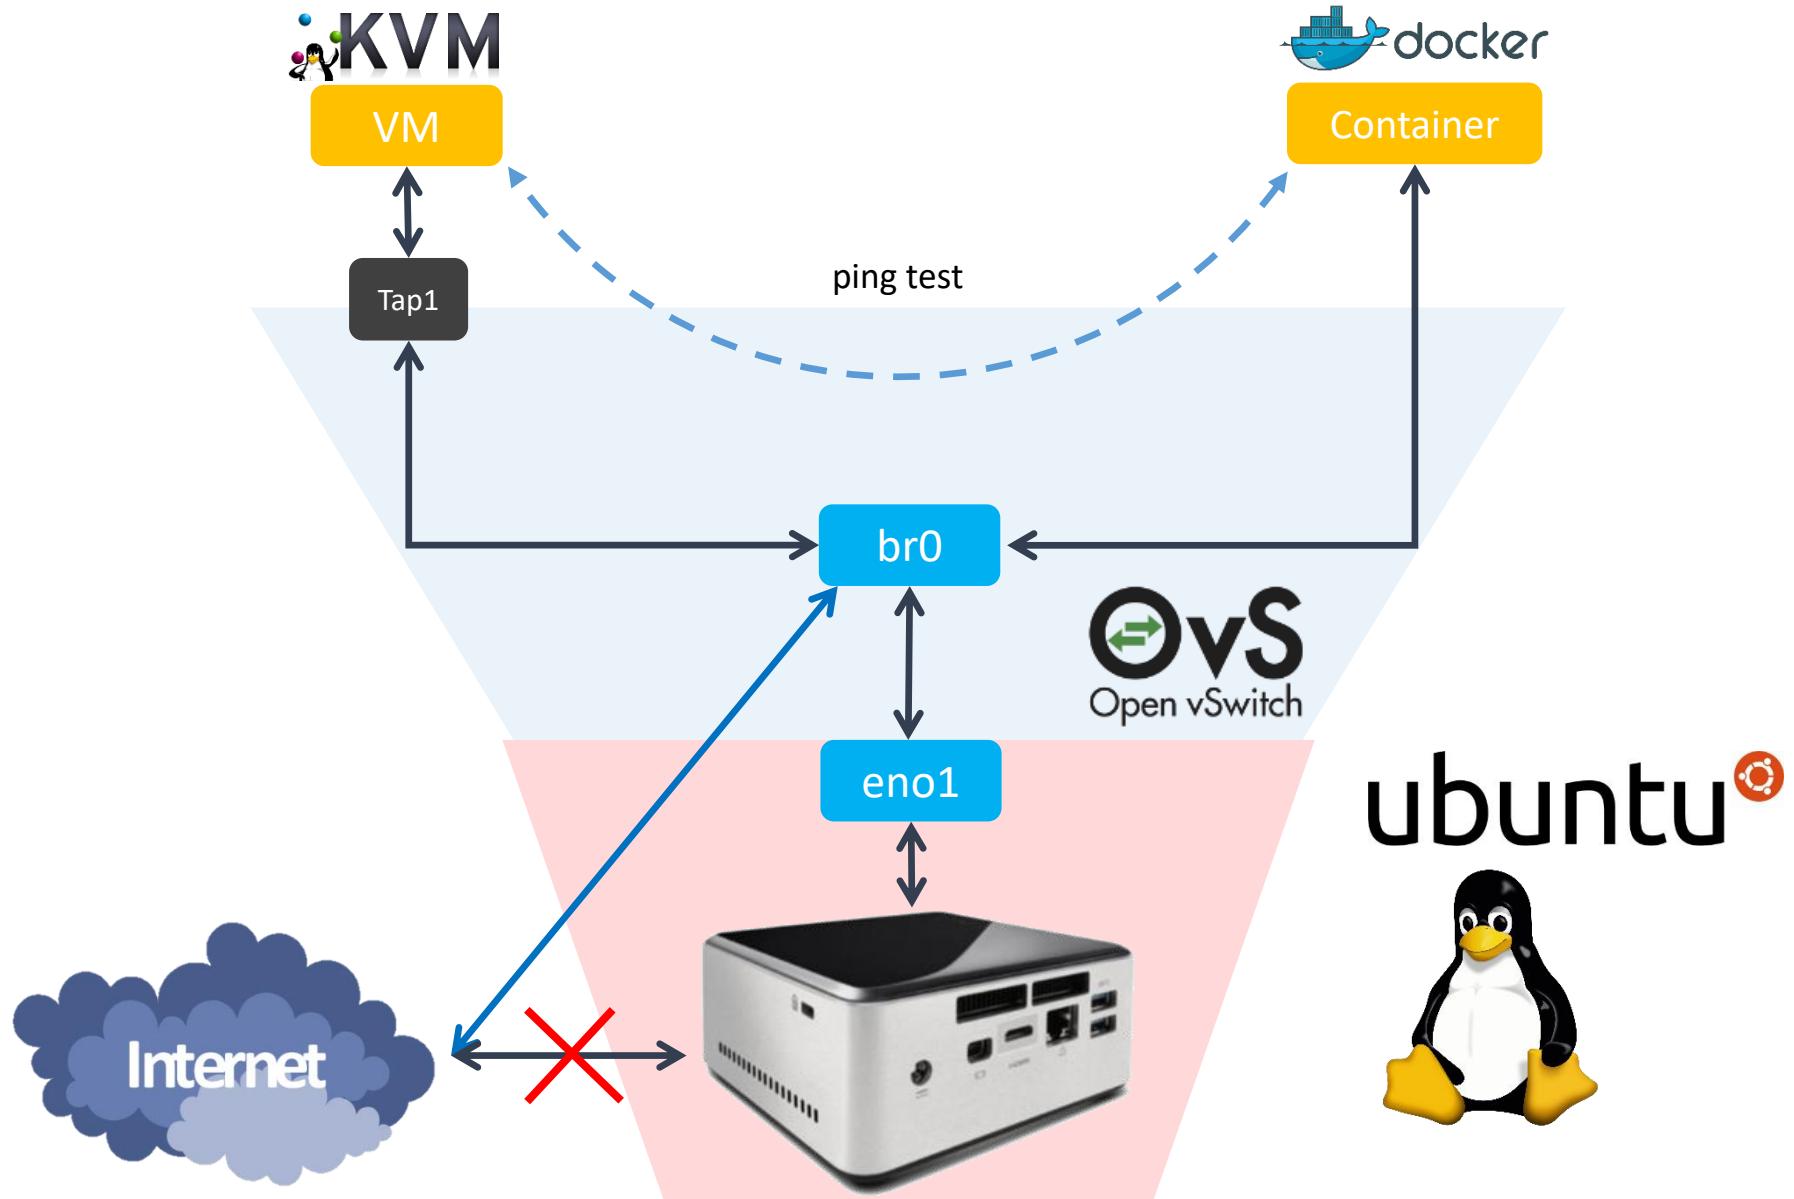
Computer Systems For AI-inspired Cloud Theory & Lab.

Lab #1: Box



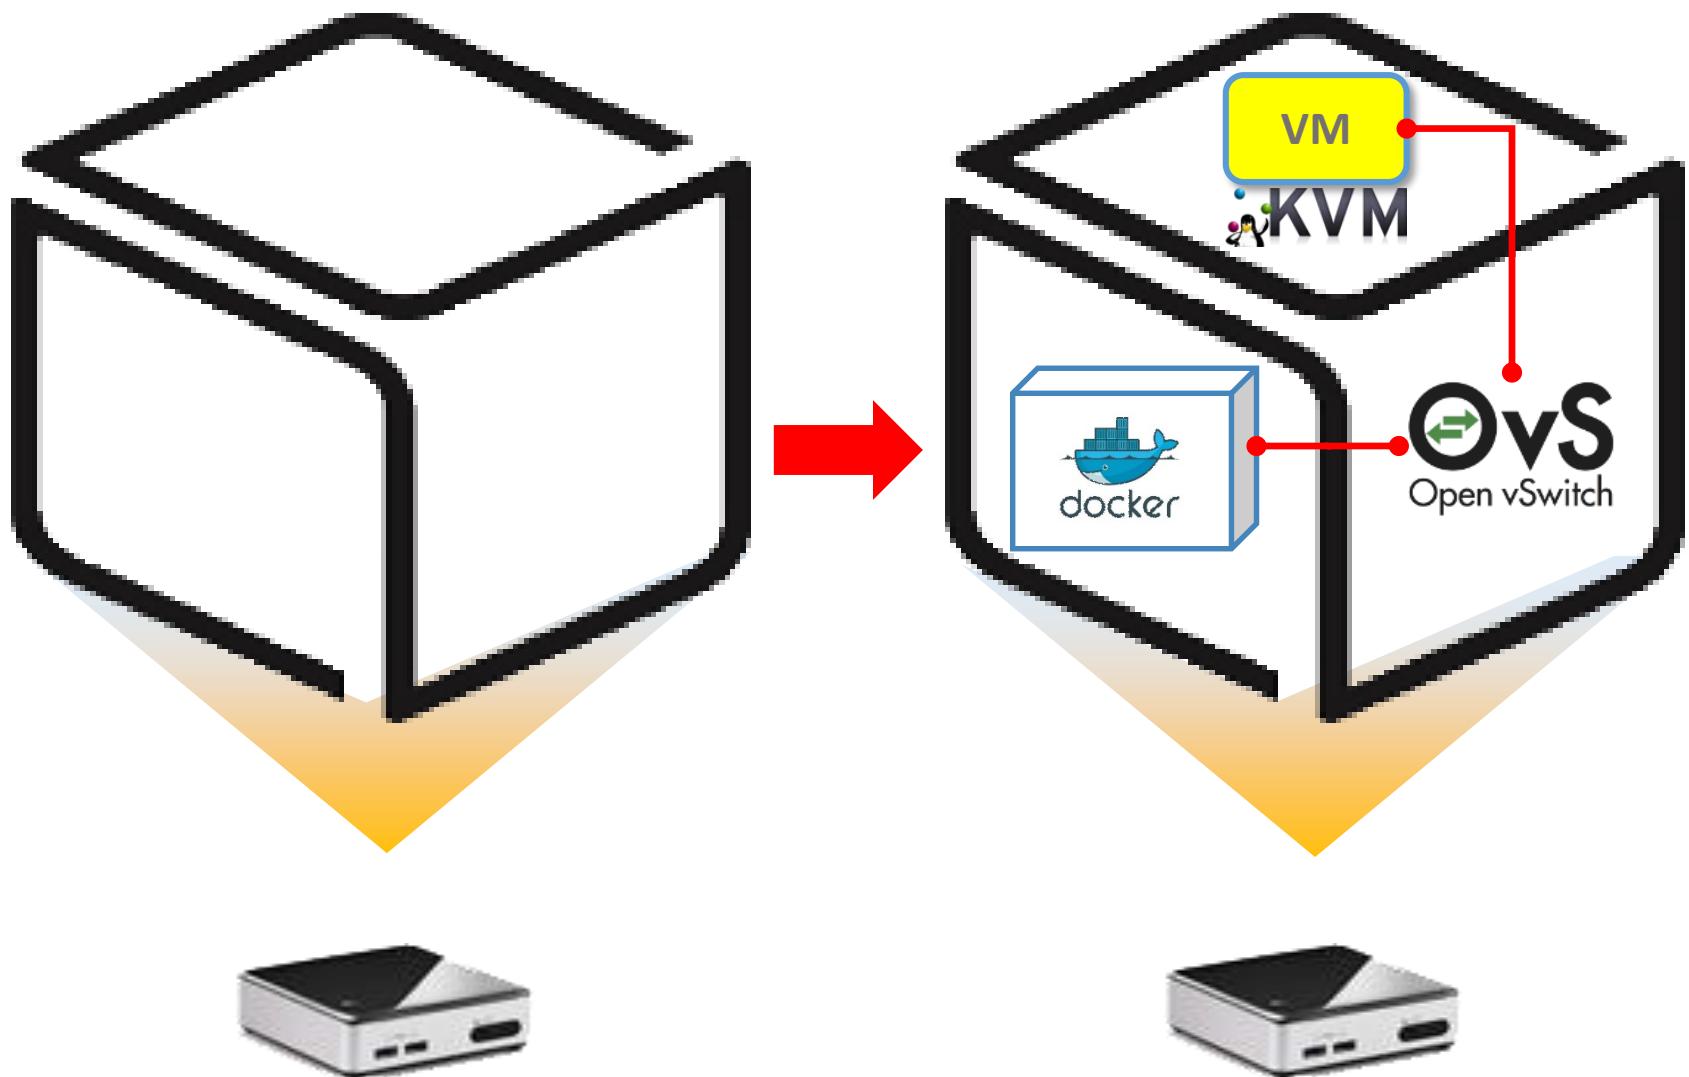
<https://github.com/SmartX-Labs/SmartX-Mini-MOOC>

Box Lab: Outline



Box Lab: Final Goal

Lab #1: Box 3



Before you start

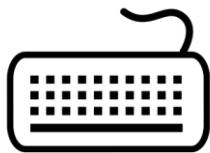
- Things you need to know

Lab
Theory

Lab
Practice

Lab
Review

Lectures are divided into Lab Theory, Lab Practice and Lab Review parts.

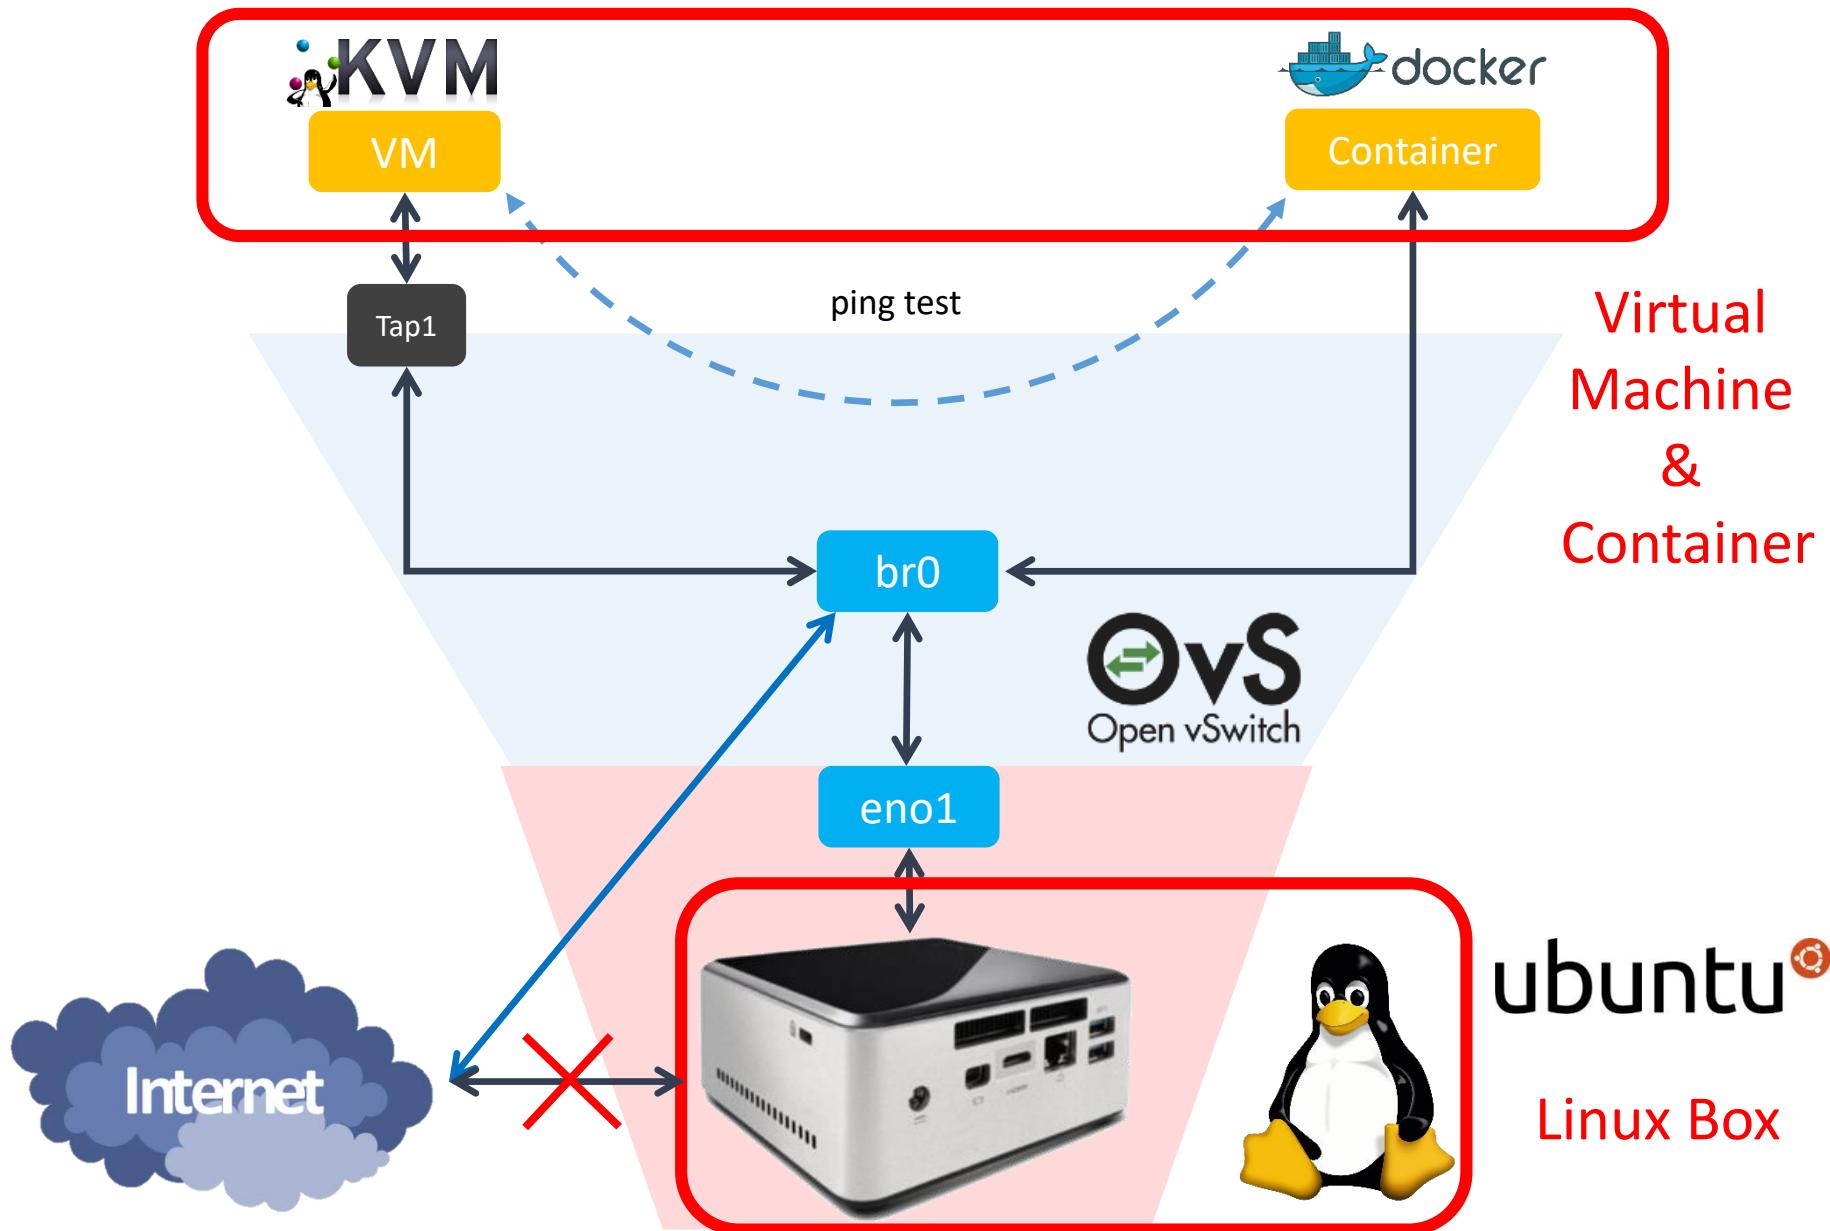


The keyboard means that you should execute instructions by following the guidance.

Lab Theory



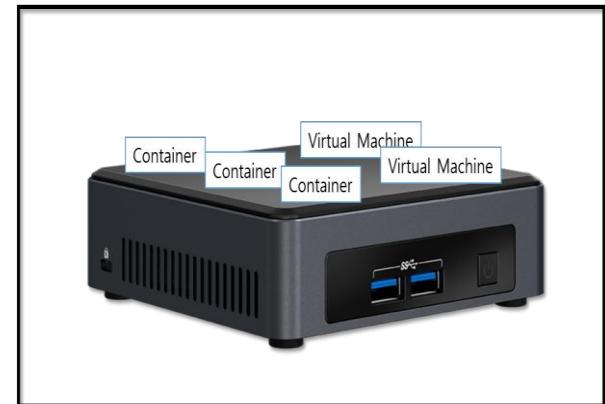
Linux Box with Virtualization/Containers



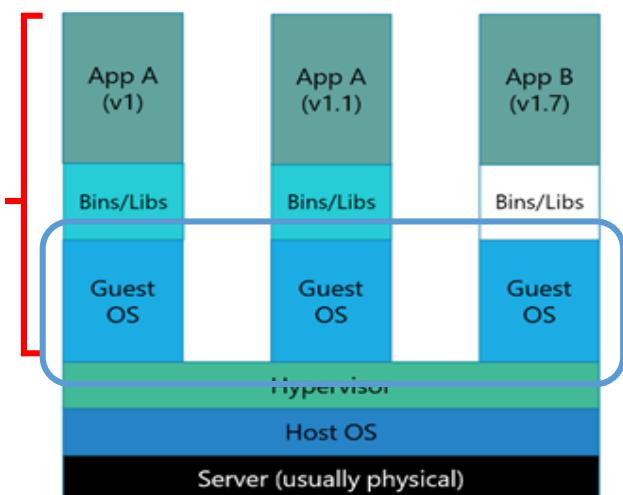
VMs and Containers on a Linux Box

On a Linux Box (e.g., NUC)

- Multiple fully-isolated (with dedicated resources), but heavy VM instances (i.e., VMs).
- Multiple partially-isolated, lightweight container instances (i.e., containers)

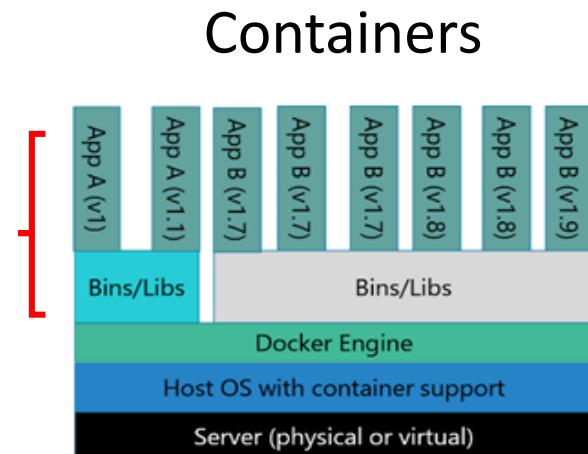


Virtual Machines (VMs)



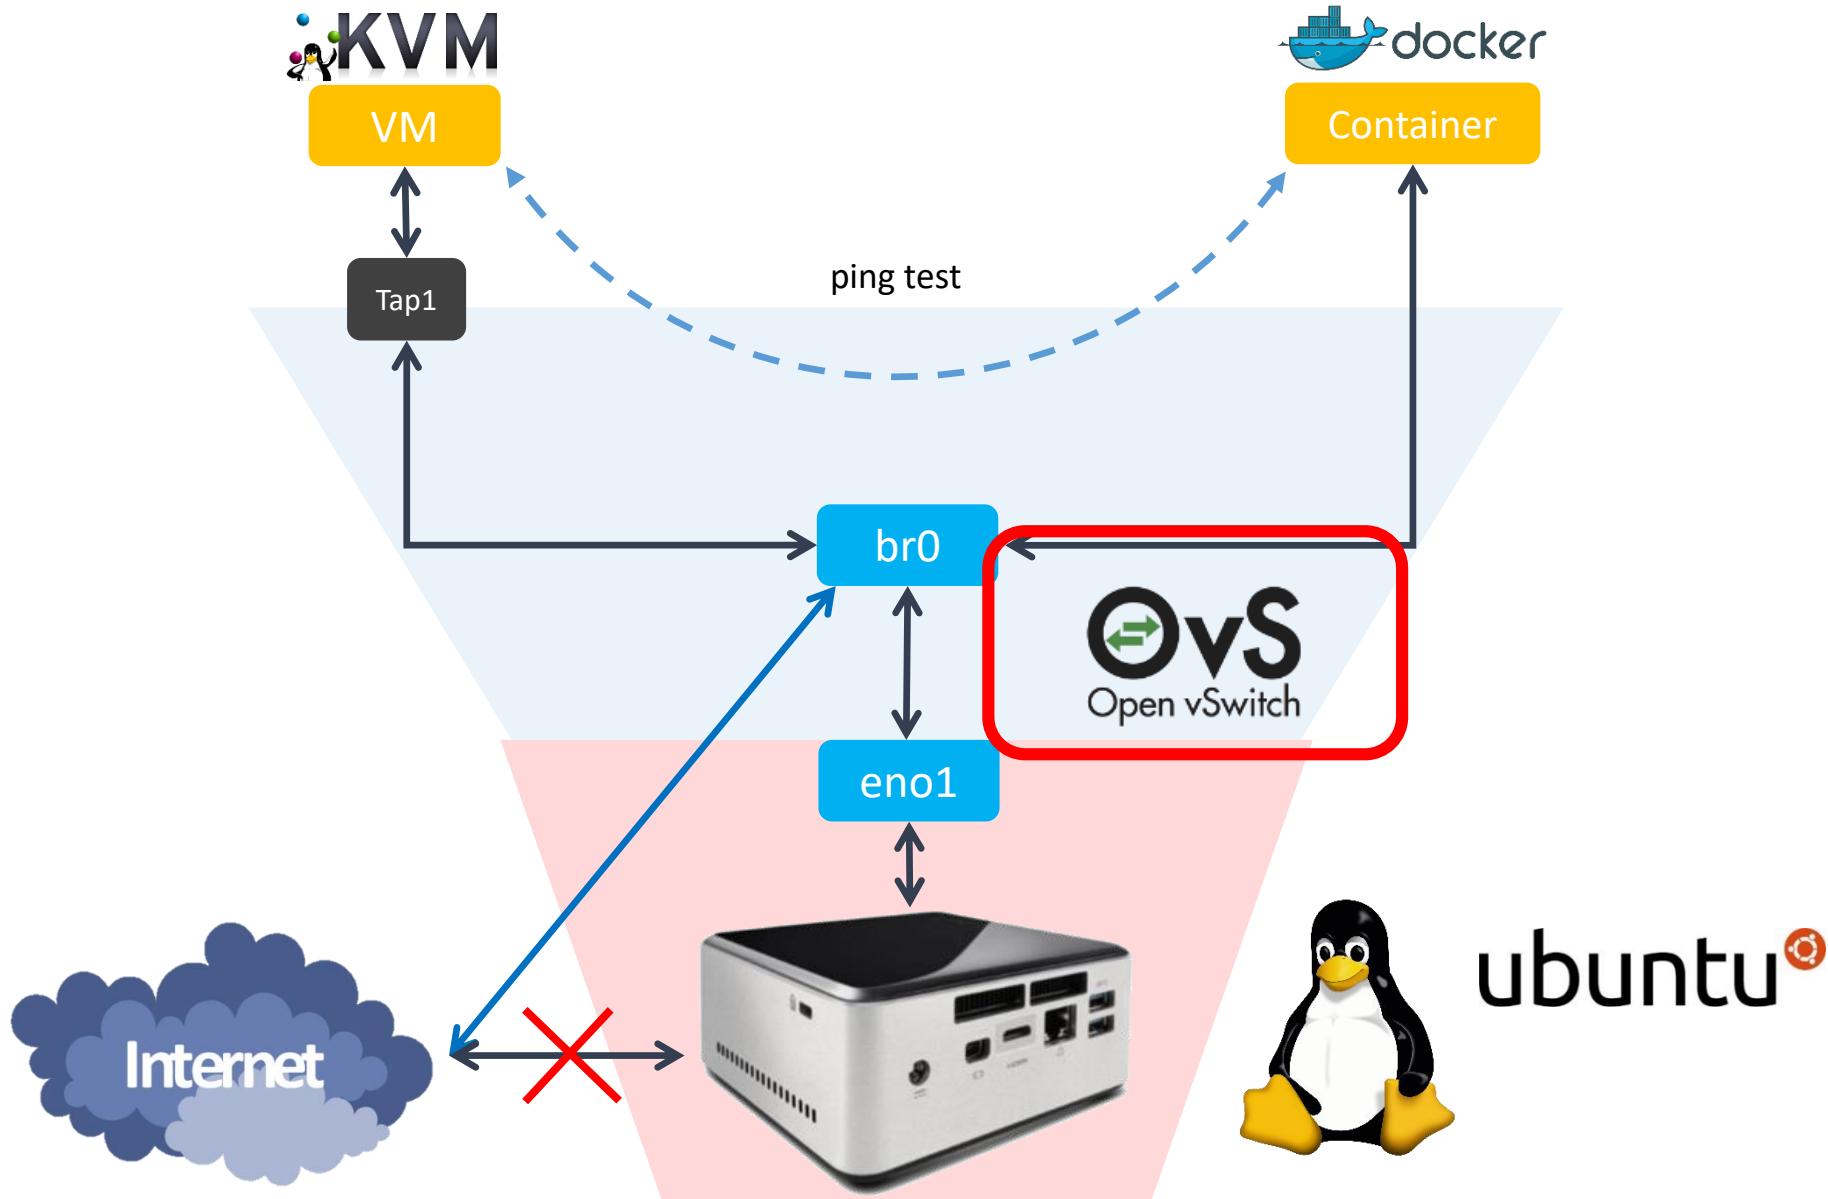
Hypervisor Tool for creating VMs

Heavier than Container!

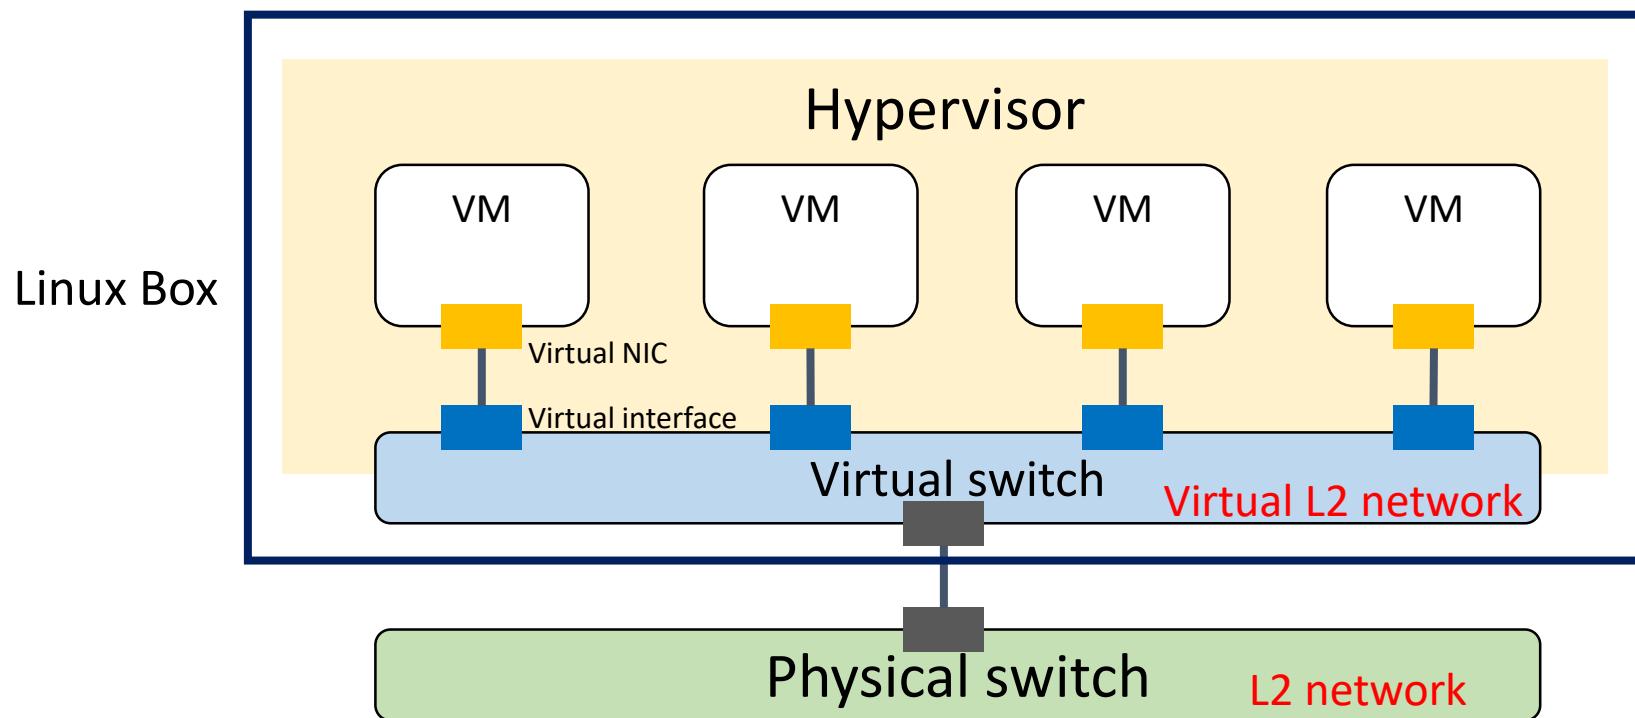


Container Runtime Tool for running containers

A Switch inside Linux Box: Open vSwitch



Virtual Switch in a Box to connect VMs



- A software-based virtual switch allows one VM to communicate with neighbor VMs as well as to connect to Internet (via physical switch).
- Software-based switches (running with the power of CPUs) are known to be more flexible/upgradable and benefited of virtualization (memory overcommit, page sharing, ...)
- VMs (similarly containers) have logical (virtual) NIC with virtual Ethernet ports so that they can be plugged into the virtual interface (port) of virtual switches.

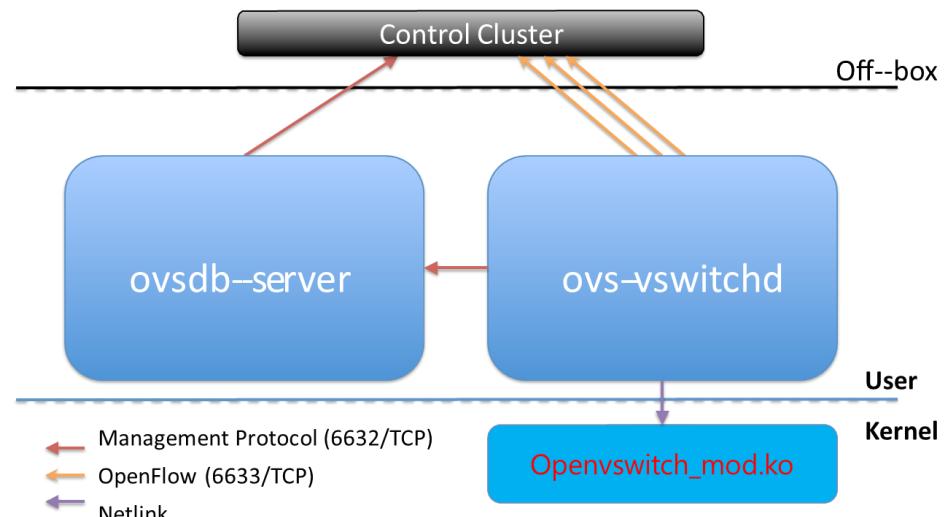
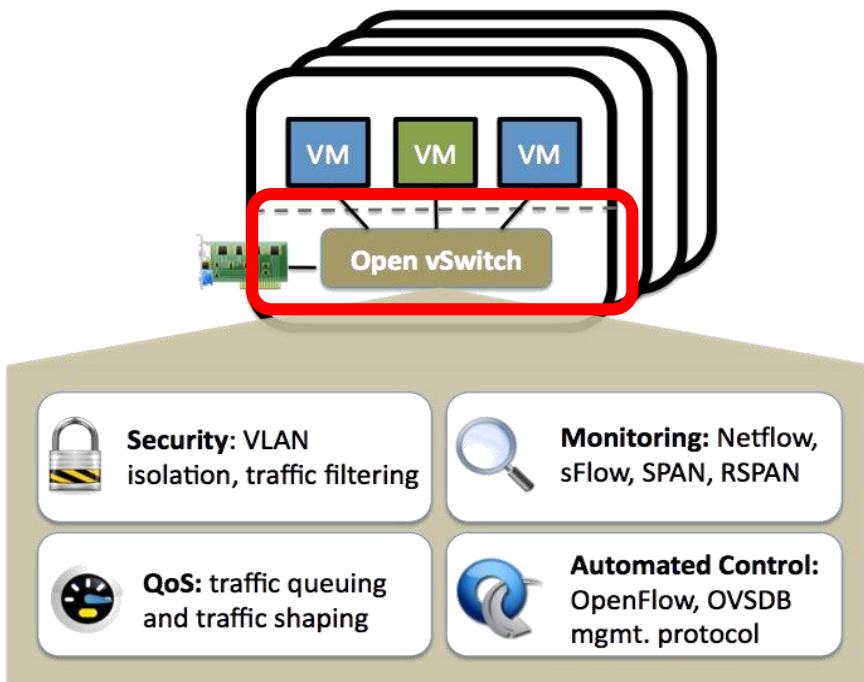
Linux-adopted virtual switch: Open vSwitch



<http://openvswitch.org/>



Open vSwitch is an open-source virtual switch software designed for virtual servers.



OVS Main components

Lab Practice



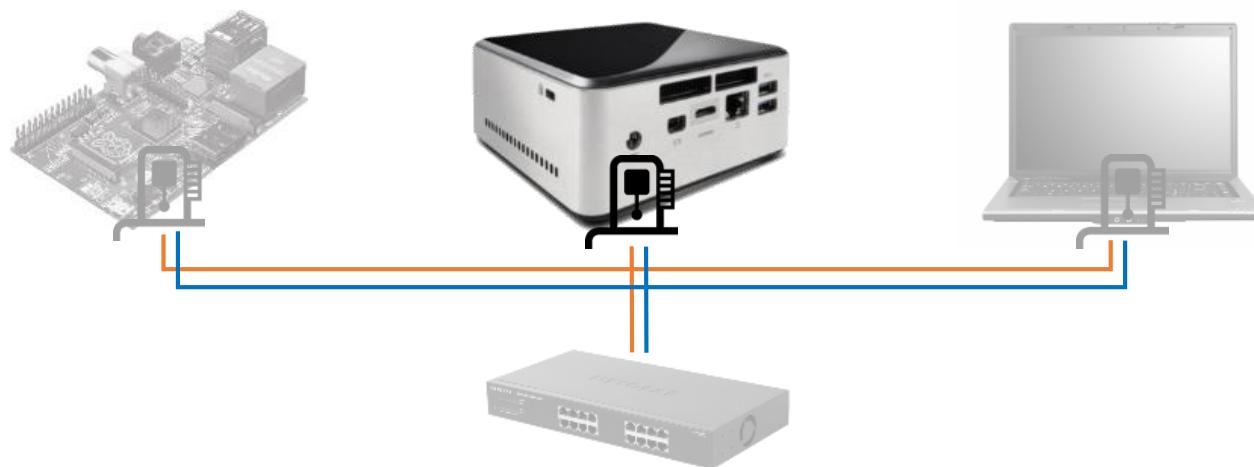
#0 - Lab Preparation

Wired connection

NAME: Raspberry Pi Model B (Pi)
CPU: ARM Cortex A7 @900MHz
CORE: 4
Memory: 1GB
SD Card: 32GB

NAME: NUC5i5MYHE (NUC PC)
CPU: Intel® Core™ i3-8109U 3.60 GHz
CORE: 2
Memory: 32GB DDR4
HDD: 94GB

NAME: NT900X3A
CPU: i5-2537U @1.40GHz
CORE: 2
Memory: 4GB DDR3
HDD: 128GB



NAME: netgear prosafe 16 port gigabit switch(Switch)
Network Ports: 16 auto-sensing 10/100/1000 Mbps Ethernet ports

#1 - NUC: OS Installation



- OS : Ubuntu Desktop 20.04 LTS(64bit)
 - Download Site : <https://releases.ubuntu.com/20.04/>

A screenshot of the Ubuntu releases website. The top navigation bar says "ubuntu releases". Below it, a large orange header box displays "Ubuntu 20.04.2.0 LTS (Focal Fossa)". Underneath, there's a section titled "Select an image" with a sub-section titled "Desktop image". It explains that the desktop image allows trying Ubuntu without changing the computer and installing it later, requiring at least 1024MB of RAM. A sub-section for "64-bit PC (AMD64) desktop image" is also shown, noting it's for AMD64 or EM64T architecture like Athlon64, Opteron, EM64T Xeon, and Core 2 processors.



Ubuntu Home
Screen After
Installation

#1 - NUC: OS Installation

- OS Installation



Updates and other software

- Select 'Minimal installation'

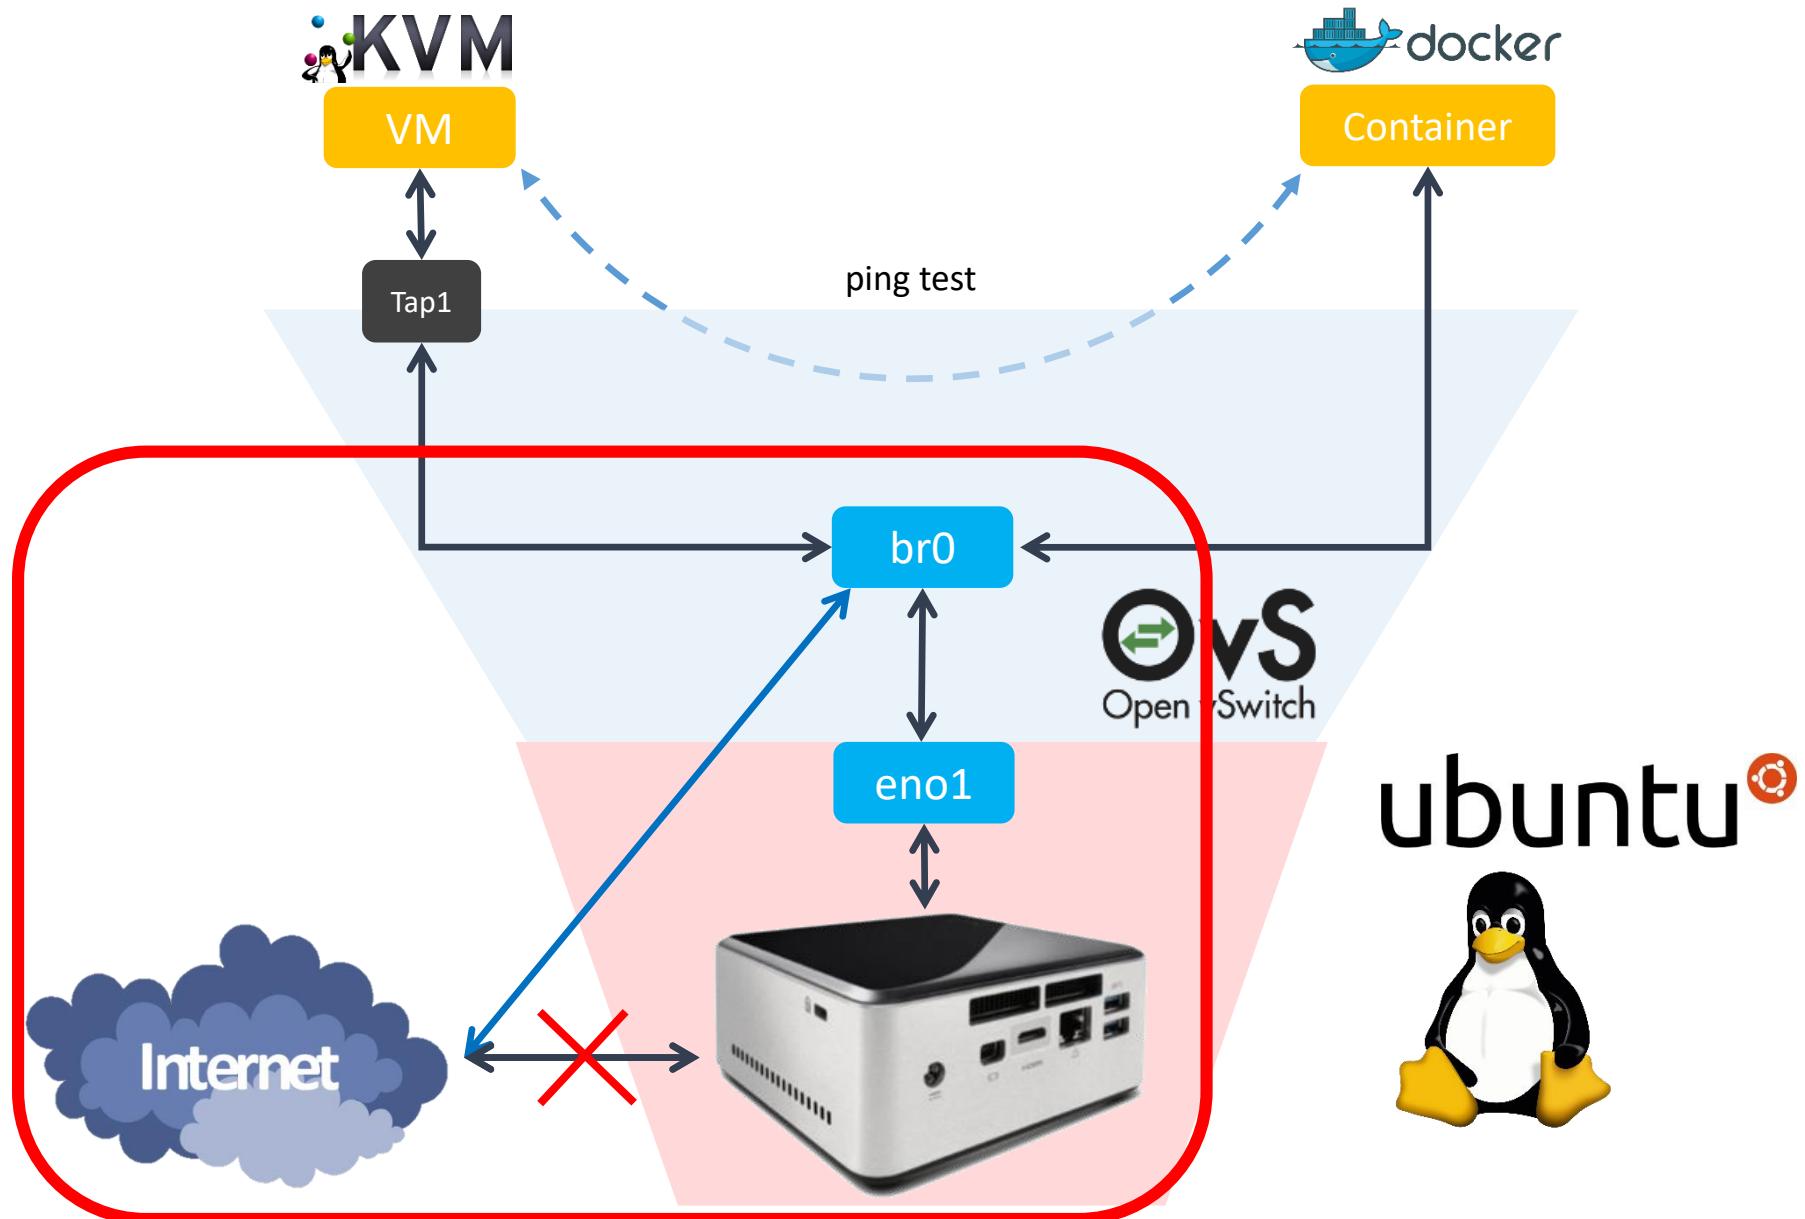
Installation type

- Select 'Something else'
- On /dev/sda or /dev/nvme0n1
 - (UEFI), add 512MB EFI partition
 - Add empty partition with 20GB (20480MB) (Select 'do not use the partition')
 - Add Etc4 partition on leave memory
- Select Boot loader
 - BIOS: Ext4 partition
 - UEFI: EFI partition

LVM 관련 오류 발생 시

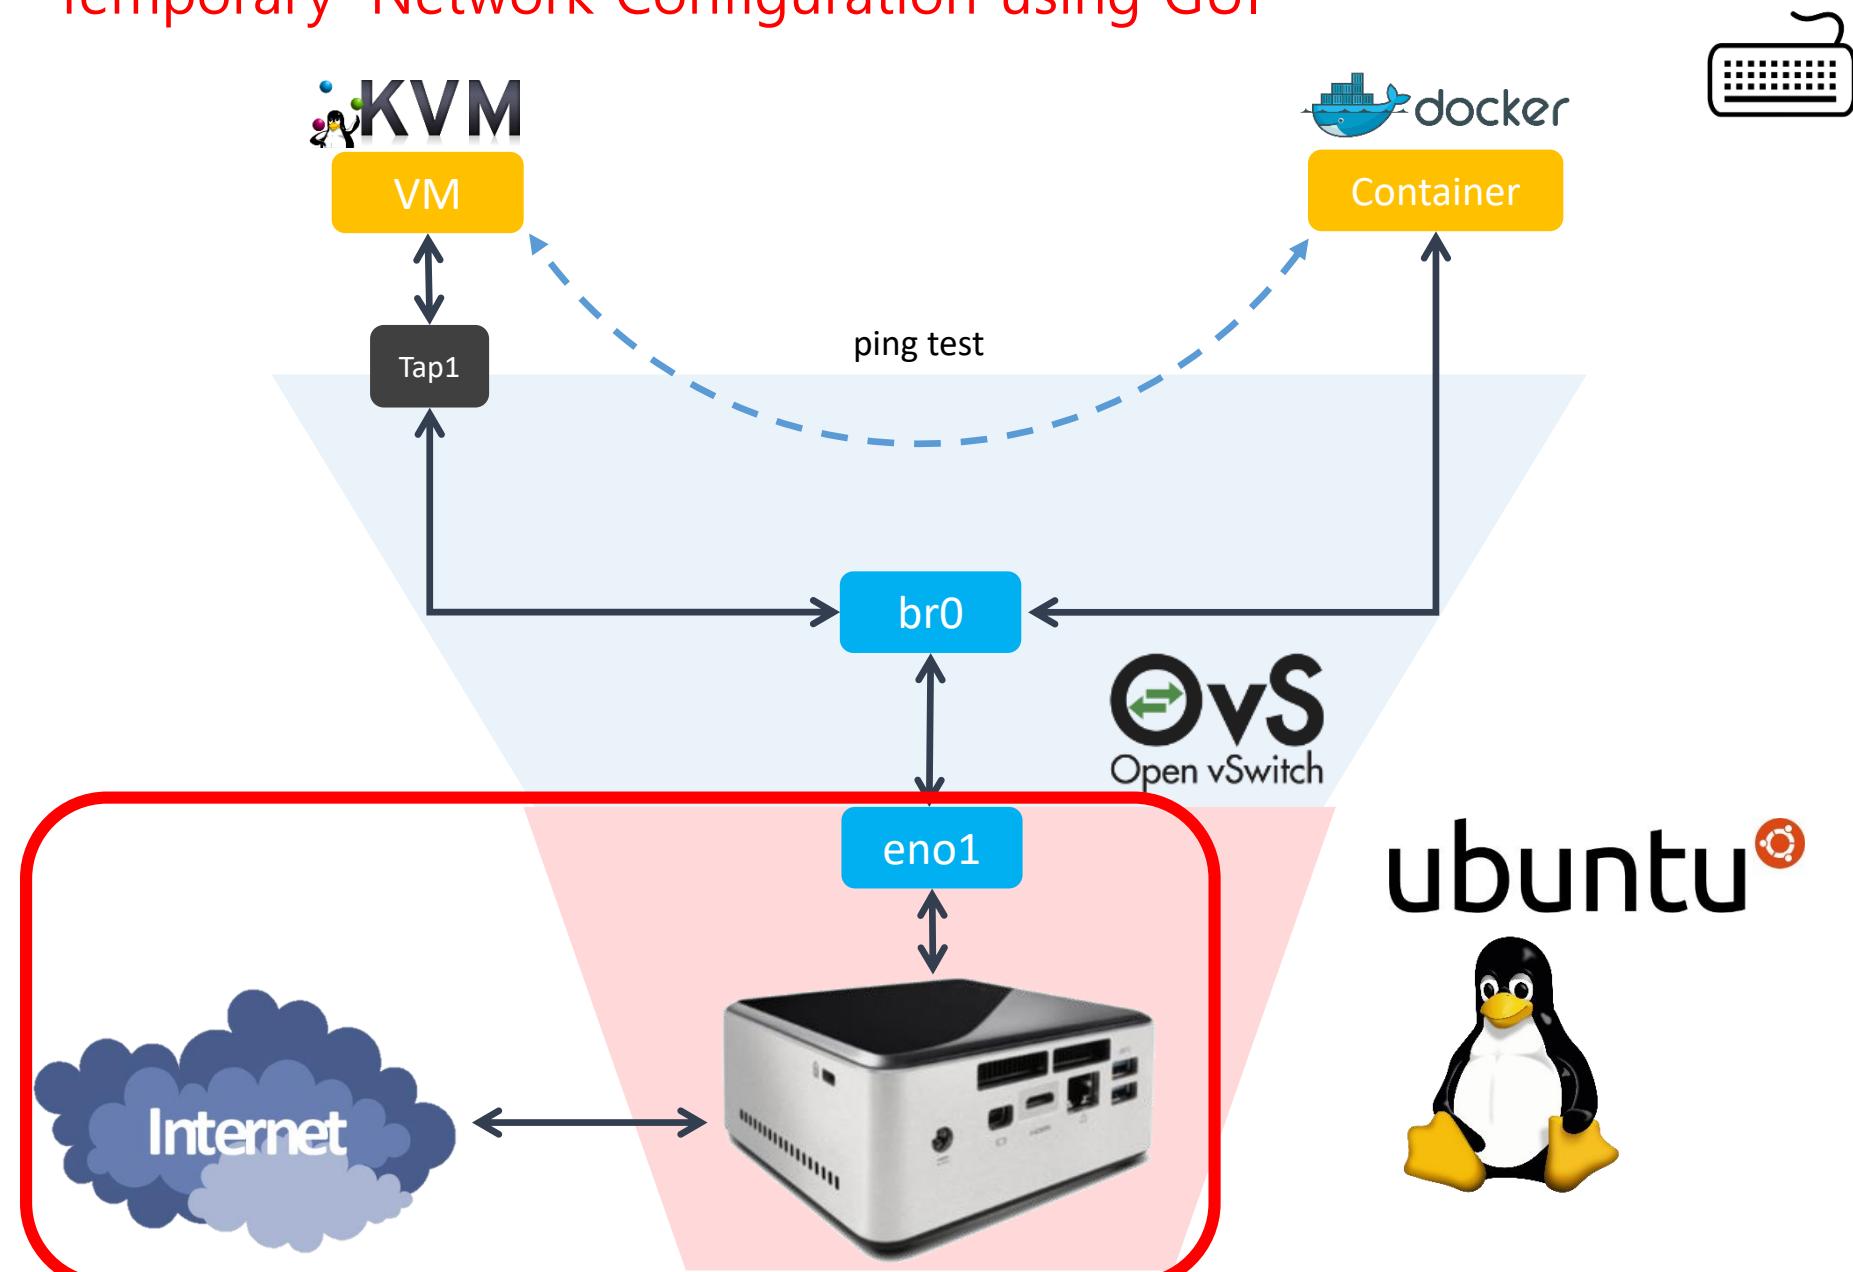
1. 뒤로 이동하여 **Installation type** 화면으로 이동
2. Erase disk 선택
3. 시간대 선택화면까지 진행
4. 다시 뒤로 돌아가서 다시 **Installation type** 화면으로 이동
5. Something else 선택 후 정상 진행

#2 - NUC: Network Configuration



#2 - NUC: Network Configuration

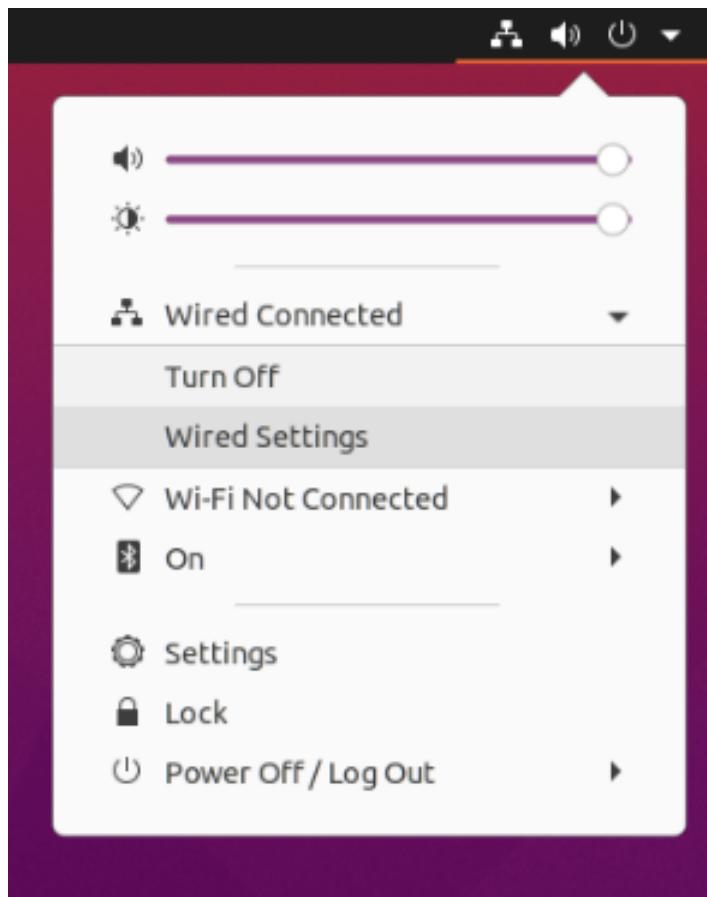
- 'Temporary' Network Configuration using GUI



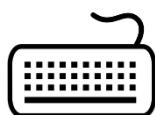
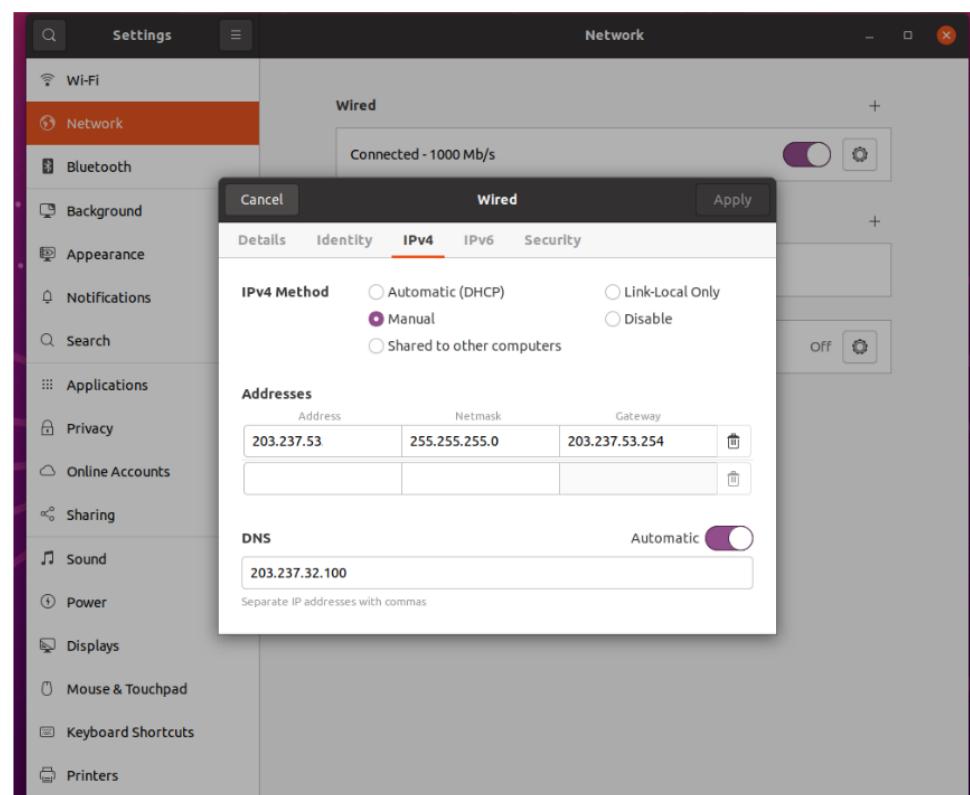
#2 - NUC: Network Configuration

- 'Temporary' Network Configuration using GUI

Click the LAN configuration icon.



*Enter the network info.
(IP address, subnet mask, gateway)*



#2 - NUC: Network Configuration

- Set Prerequisites



0. Update & Upgrade

```
$sudo apt update  
$sudo apt upgrade
```

1. Install net-tools & ifupdown

```
$sudo apt install net-tools ifupdown  
$ifconfig -a
```

```
netcs@netcs-desktop:~$ ifconfig -a  
eno1: flags=4163<UP,BROADCAST,RUNNING,MULTICAST> mtu 1500  
        inet 192.168.1.199 brd 192.168.1.255 netmask 255.255.255.0 broadcast 192.168.1.255  
        inet6 fe80::82bb:956:11d6:e640 brd fe80::ff:fe11d6:e640/64 scopeid 0x20<link>  
          ether ec:a8:6b:fb:a2:09 txqueuelen 1000 (Ethernet)  
            RX packets 255746 bytes 382878532 (382.8 MB)  
            RX errors 0 dropped 0 overruns 0 frame 0  
            TX packets 40500 bytes 3146671 (3.1 MB)  
            TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0  
            device interrupt 20 memory 0xf7c00000-f7c20000  
  
lo: flags=73<UP,LOOPBACK,RUNNING> mtu 65536  
        inet 127.0.0.1 brd 127.0.0.1 netmask 255.0.0.0  
        inet6 ::1 brd :: prefixlen 128 scopeid 0x10<host>  
          loop txqueuelen 1000 (Local Loopback)  
            RX packets 2771 bytes 237705 (237.7 KB)  
            RX errors 0 dropped 0 overruns 0 frame 0  
            TX packets 2771 bytes 237705 (237.7 KB)  
            TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0
```

#2 - NUC: Network Configuration

- Set Prerequisites



2. Install openvswitch-switch & make br0 bridge

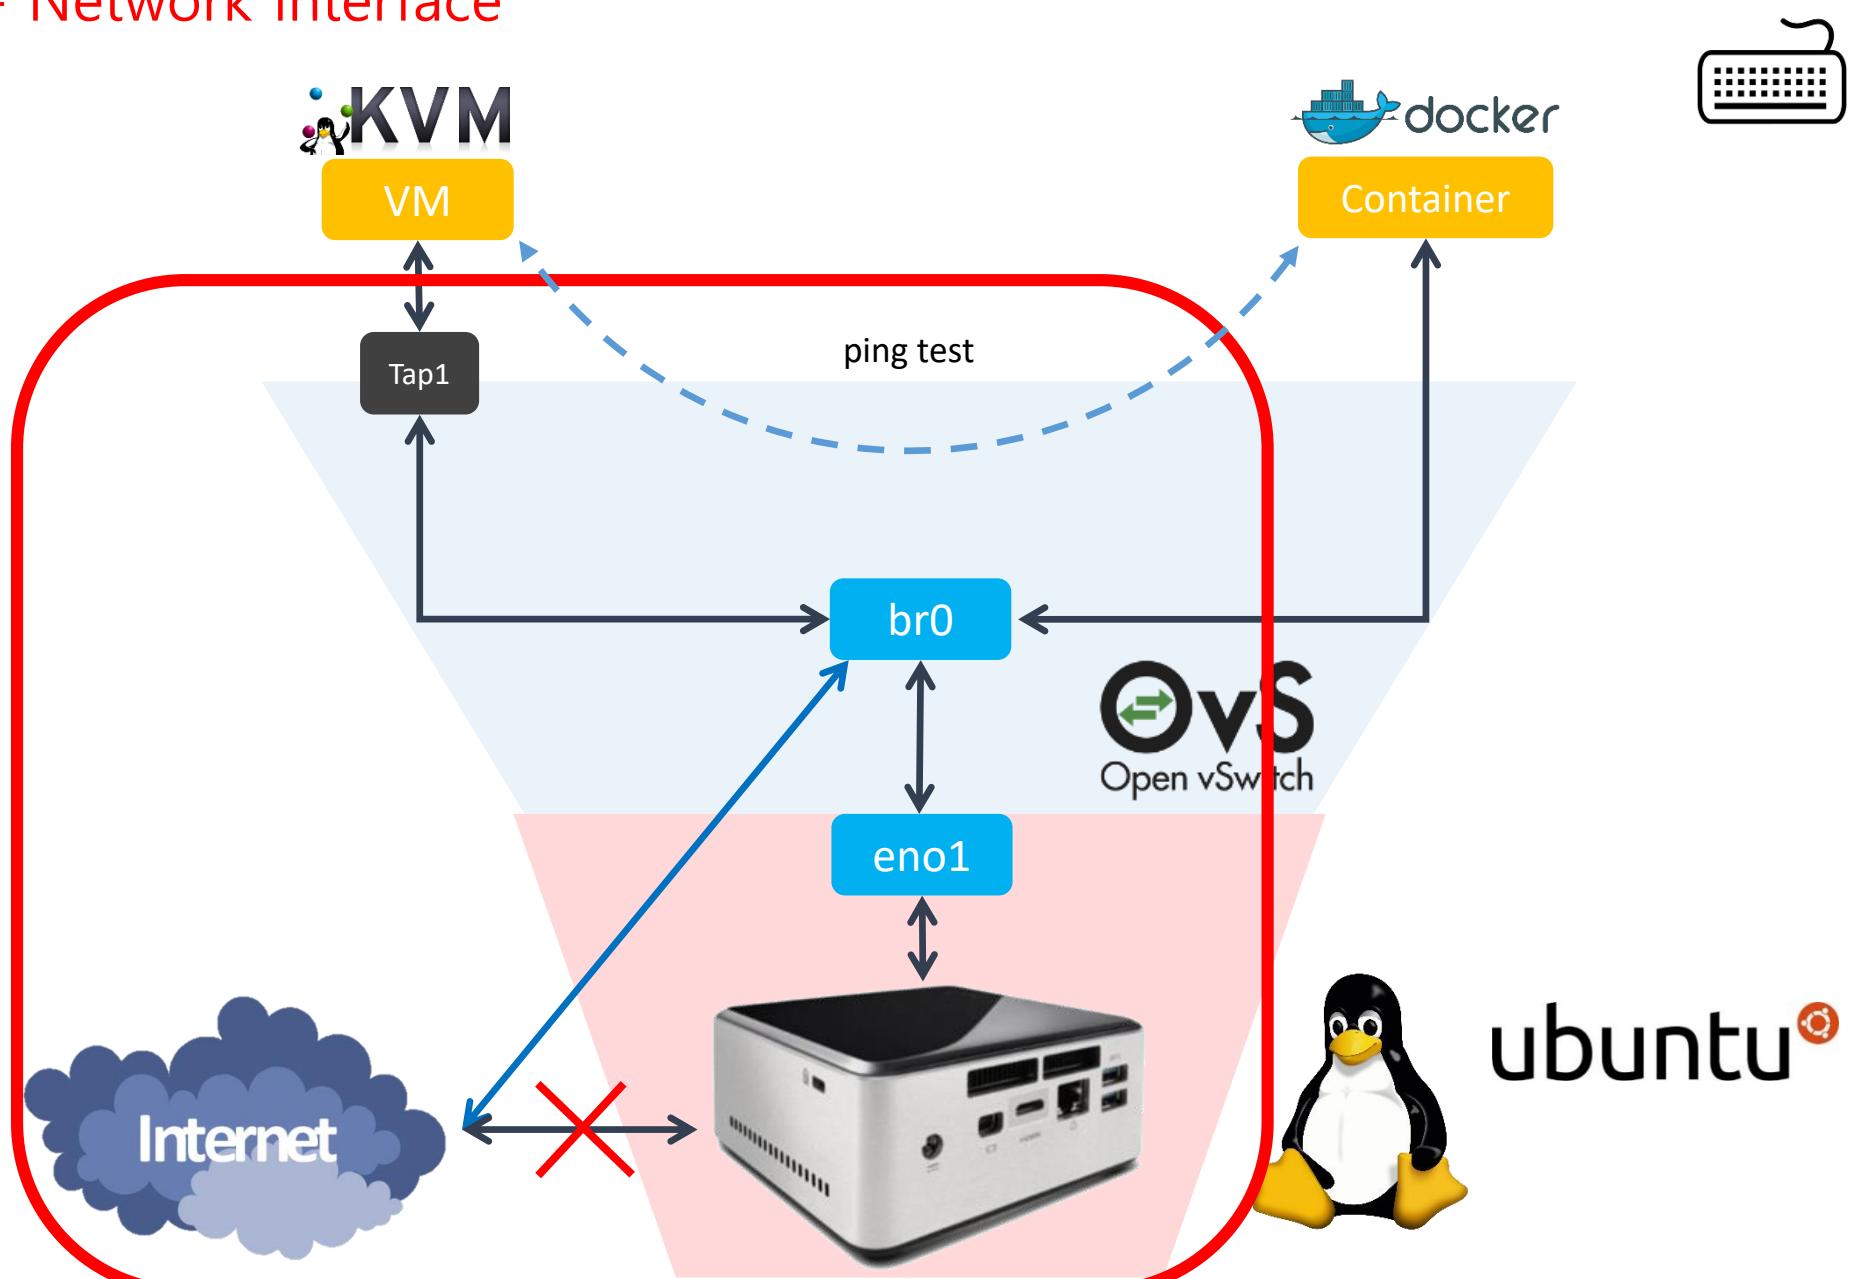
```
$sudo apt intall openvswitch-switch  
$sudo ovs-vsctl add-br br0  
$sudo ovs-vsctl show
```

```
nuc1@nuc1-NUC8i3BEK:~$ sudo ovs-vsctl show  
341b772b-b8b8-45b5-8a07-3fc1a6ebb2d9  
    Bridge br0  
        Port br0  
            Interface br0  
                type: internal  
        ovs_version: "2.13.1"  
nuc1@nuc1-NUC8i3BEK:~$ █
```

ovs-vsctl show

#2 - NUC: Network Configuration

- Network interface



#2 - NUC: Network Configuration

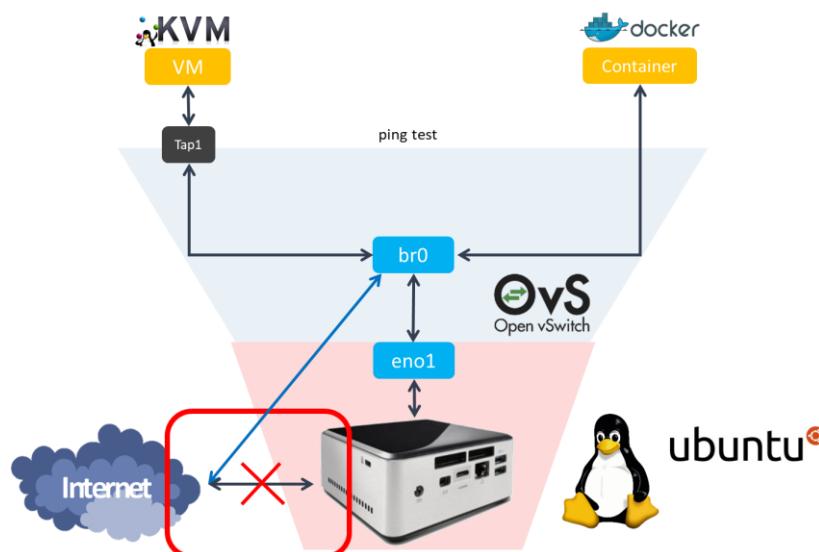
- Network interface

Disable netplan



```
$sudo su
```

```
#systemctl stop systemd-networkd.socket systemd-networkd networkd-dispatcher systemd-networkd-wait-online  
#systemctl disable systemd-networkd.socket systemd-networkd networkd-dispatcher systemd-networkd-wait-online  
#systemctl mask systemd-networkd.socket systemd-networkd networkd-dispatcher systemd-networkd-wait-online  
#apt-get --assume-yes purge nplan netplan.io  
#exit
```



#2 - NUC: Network Configuration

- Network interface

Open /etc/network/interfaces via 'vi'



```
$sudo vi /etc/network/interfaces
```

Configure the Network Interface vport_vFunction is a tap interface and attach it to your VM.

----- /etc/network/interfaces -----

```
auto lo
iface lo inet loopback
```

```
auto br0
iface br0 inet static
    address ---your nuc ip---
    netmask 255.255.255.0
    gateway ---gateway ip---
    dns-nameservers 8.8.8.8
```

```
auto eno1
iface eno1 inet manual
```

```
auto vport_vFunction
iface vport_vFunction inet manual
    pre-up ip tunctl add vport_vFunction mode tap
    up ip link set dev vport_vFunction up
    post-down ip link del dev vport_vFunction
```

Some NUC have different
Interface name.
So you need to check your
NUC's interface name using
'ip a' command.

```
nuc1@nuc1-NUC8i3BEK: ~
auto lo
iface lo inet loopback

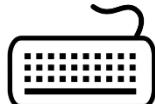
auto br0
iface br0 inet static
    address 203.237.■■■■■
    netmask 255.255.255.0
    gateway 203.237.53.254
    dns-nameservers 203.237.32.100

auto eno1
iface eno1 inet manual

auto vport_vFunction
iface vport_vFunction inet manual
    pre-up ip tunctl add vport_vFunction mode tap
    up ip link set dev vport_vFunction up
    post-down ip link del dev vport_vFunction
```

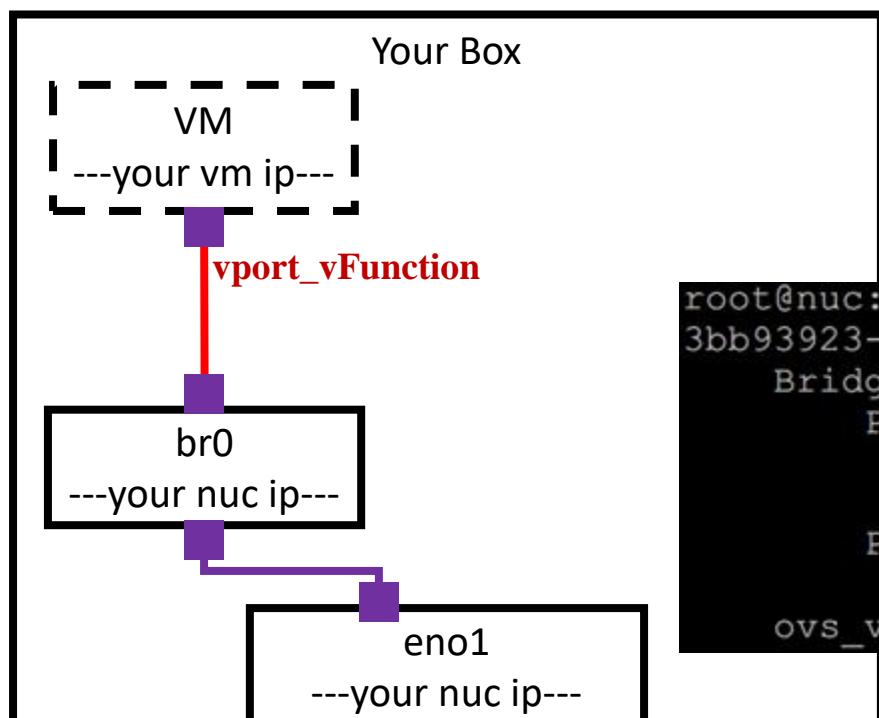
#2 - NUC: Network Configuration

- Network interface



We will make VM attaching vport_vFunction. You can think this tap as a NIC of VM.

Below is the figure you configurated so far



```
root@nuc:~# ovs-vsctl show
3bb93923-3eac-420a-9da9-9143aff14209
    Bridge "br0"
        Port "br0"
            Interface "br0"
                type: internal
        Port vport_vFunction → [tap_name]
            Interface vport_vFunction
    ovs_version: "2.0.2"
```

#2 - NUC: Network Configuration

- Network interface

Restart the whole interfaces 1



```
$sudo su  
#systemctl unmask networking  
#systemctl enable networking  
#systemctl restart networking
```

add port 'eno1' and 'vport_vFunction' to 'br0'

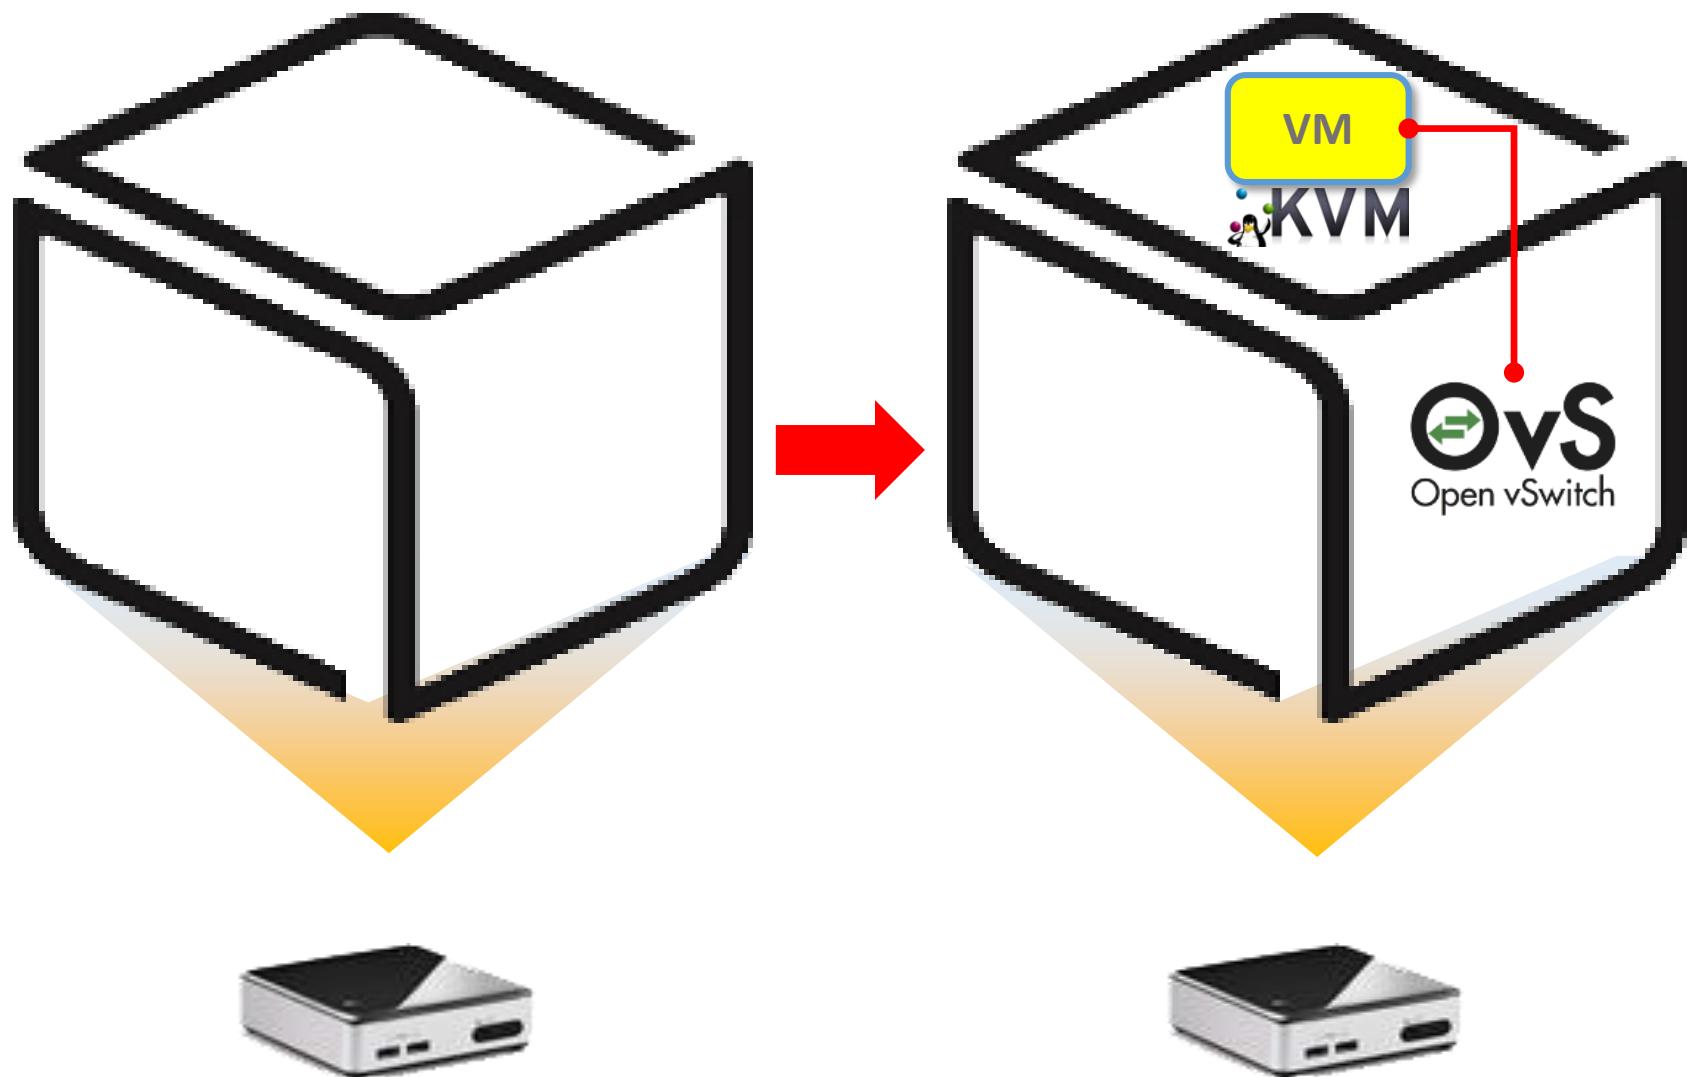
```
$ sudo ovs-vsctl add-port br0 eno1  
$ sudo ovs-vsctl add-port br0 vport_vFunction  
$ sudo ovs-vsctl show
```

Restart the whole interfaces 2

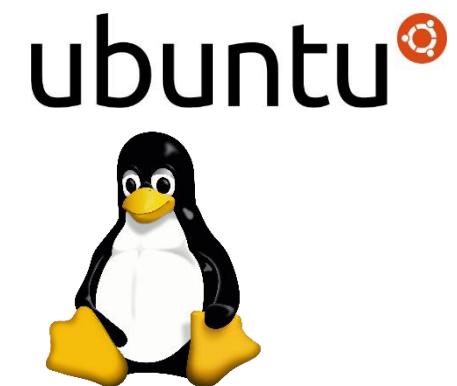
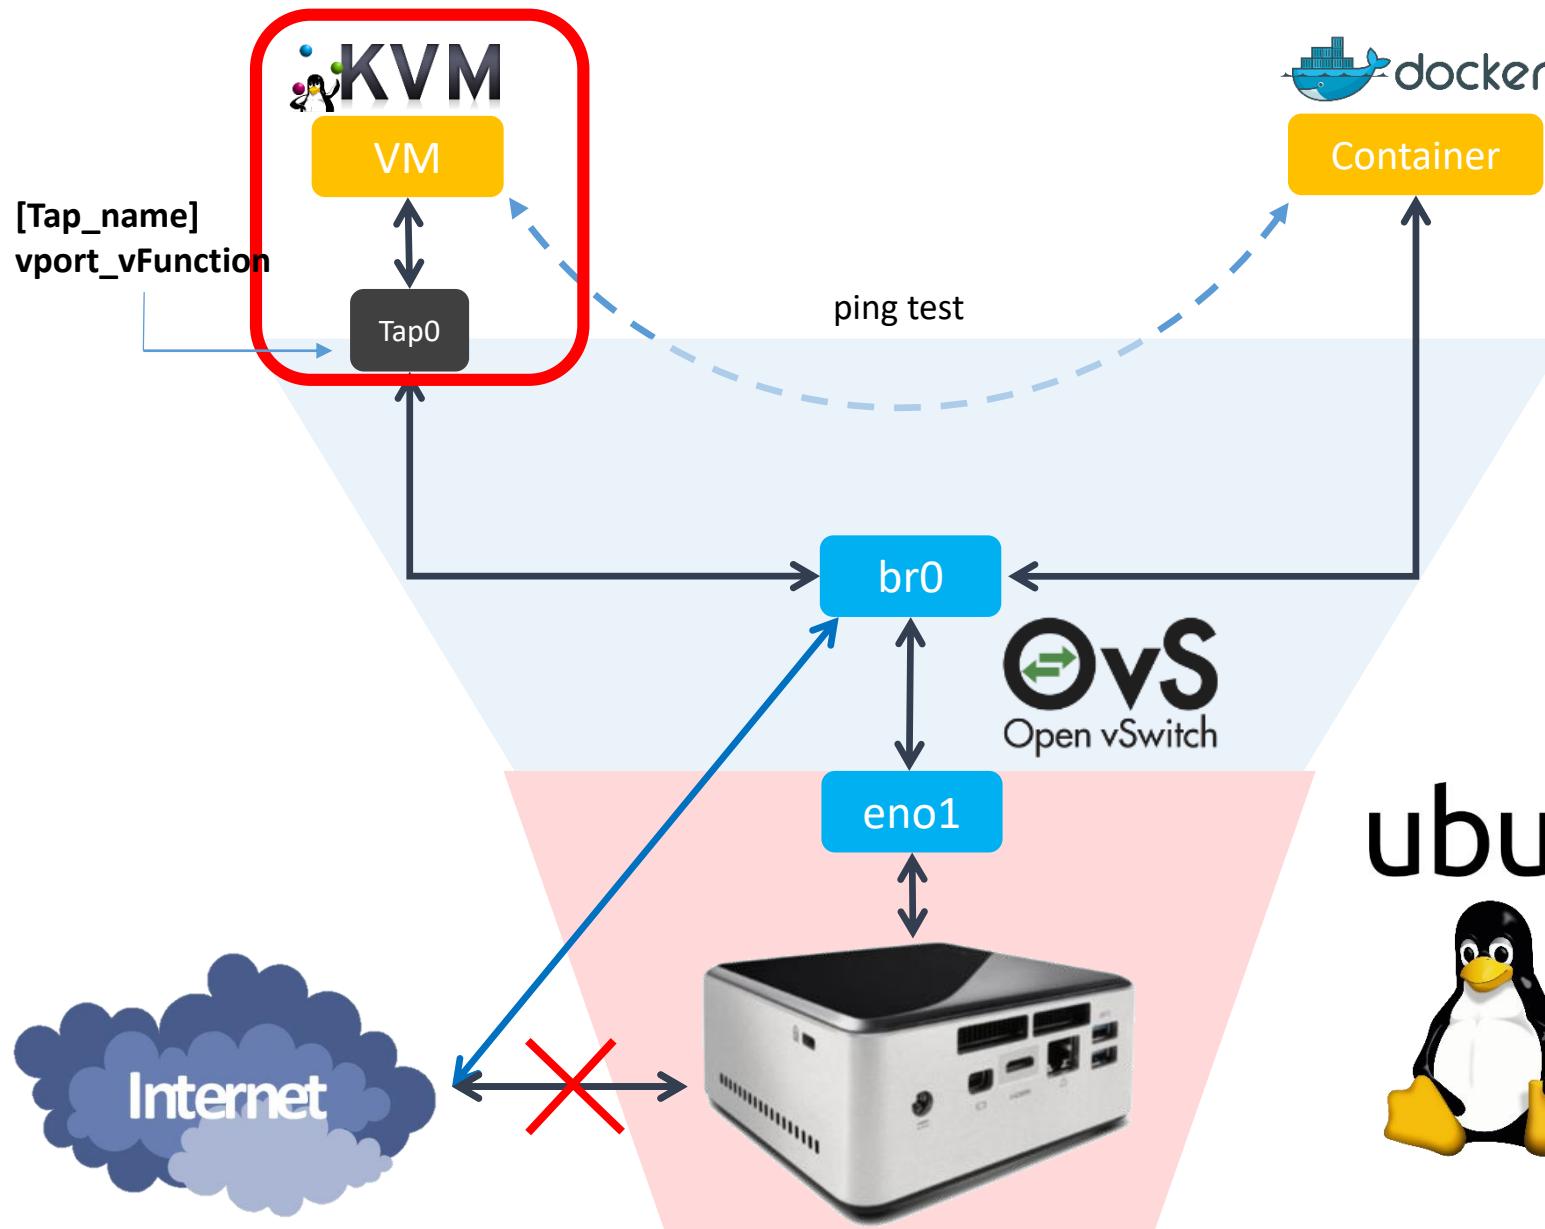
```
#systemctl unmask networking  
#systemctl enable networking  
#systemctl restart networking  
#exit
```

KVM-based VM connected via OVS

- Goal of this section



#3 - NUC: Making VM with KVM



#3 - NUC: Making VM with KVM

-Install dependency to upgrade KVM



Install dependency & download Ubuntu 20.04 64bit server image.

```
$sudo apt install qemu-kvm libvirt-daemon-system libvirt-clients bridge-utils  
//upgrade KVM  
//qemu is open-source emulator
```

```
$wget http://old-releases.ubuntu.com/releases/focal/ubuntu-20.04-beta-live-server-amd64.iso
```

Now we are ready to make VM. So continue the setting.

#3 - NUC: Making VM with KVM

-Prepare for Ubuntu VM



Make a VM image.

```
$sudo qemu-img create [img_name].img -f qcow2 [storage_capacity]
```

```
$sudo qemu-img create vFunction20.img -f qcow2 10G
```

Boot VM image from Ubuntu iso file (**띠어쓰기 주의!**).

```
$sudo kvm -m [memory_capacity] -name [vm_name] -smp cpus=[#cpu],maxcpus= [#maxcpu] -device virtio-net-pci,netdev=net0 -netdev tap,id=net0,ifname= [tap_name],script=no -boot d [img_name].img -cdrom ubuntu-20.04-beta-live-server-amd64.iso -vnc :[#] --daemonize -monitor telnet:127.0.0.1:3010,server,nowait,ipv4
```

```
$ sudo kvm -m 1024 -name tt -smp cpus=2,maxcpus=2 -device virtio-net-pci,netdev=net0 -netdev tap,id=net0,ifname=vport_vFunction,script=no -boot d vFunction20.img -cdrom ubuntu-20.04-beta-live-server-amd64.iso -vnc :5 -daemonize -monitor telnet:127.0.0.1:3010,server,nowait,ipv4
```

#3 - NUC: Making VM with KVM

-Prepare for Ubuntu VM

Configure SNAT with iptables for VM network

```
$sudo iptables -A FORWARD -i eno1 -j ACCEPT
$sudo iptables -A FORWARD -o eno1 -j ACCEPT
$sudo iptables -t nat -A POSTROUTING -s 192.168.100.0/24 -o eno1 -j SNAT --to <Your ip address>
```

\$vi /etc/sysctl.conf

#net.ipv4.ip_forward=1



net.ipv4.ip_forward=1

\$sysctl -p

→ net.ipv4.ip_forward = 1

Configuration complete

```
# /etc/sysctl.conf - Configuration file for setting system variables
# See /etc/sysctl.d/ for additional system variables.
# See sysctl.conf (5) for information.
#
#kernel.domainname = example.com

# Uncomment the following to stop low-level messages on console
#kernel.printk = 3 4 1 3

#####
# Functions previously found in netbase
#

# Uncomment the next two lines to enable Spoof protection (reverse-path filter)
# Turn on Source Address Verification in all interfaces to
# prevent some spoofing attacks
#net.ipv4.conf.default.rp_filter=1
#net.ipv4.conf.all.rp_filter=1

# Uncomment the next line to enable TCP/IP SYN cookies
# See http://lwn.net/Articles/277146/
# Note: This may impact IPv6 TCP sessions too
#net.ipv4.tcp_syncookies=1

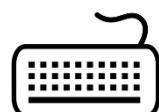
# Uncomment the next line to enable packet forwarding for IPv4
net.ipv4.ip_forward=1

# Uncomment the next line to enable packet forwarding for IPv6
# Enabling this option disables Stateless Address Autoconfiguration
# based on Router Advertisements for this host
#net.ipv6.conf.all.forwarding=1
```

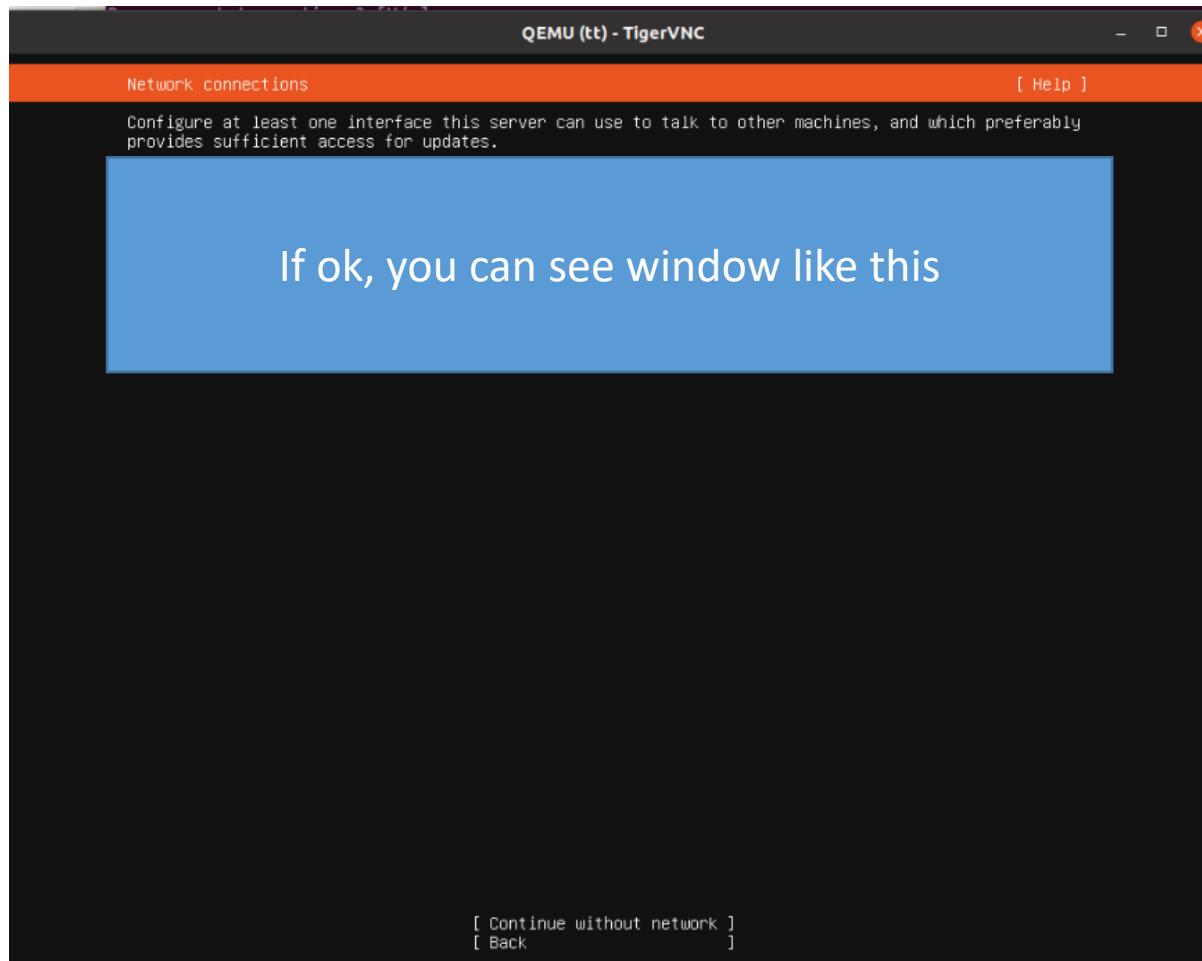
#3 - NUC: Making VM with KVM

-Install Ubuntu VM (control with 'Enter key' and 'Arrow keys')

Install VNC viewer and see inside of VM



```
$ sudo apt-get install tigervnc-viewer  
$ vncviewer localhost :5
```





#3 - NUC: Making VM with KVM

- VM network configuration (control with 'Enter key' and 'Arrow keys')

1. 네트워크 디바이스 선택 -> Edit IPv4
2. IPv4 Method -> Manual
3. 설정

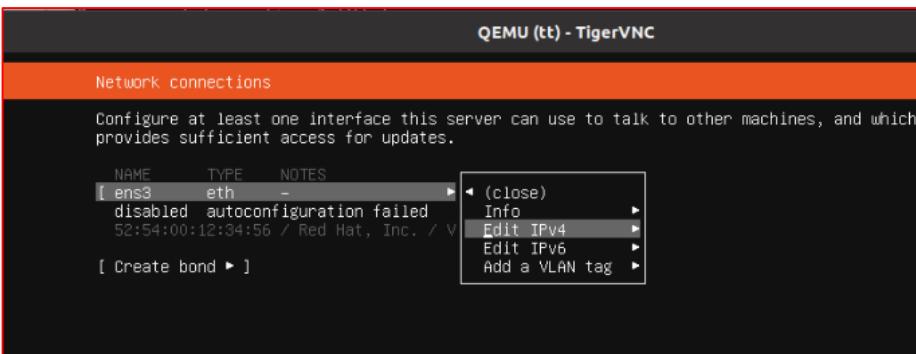
subnet: 203.237.53.0/24

Address: **---your vm ip---**

Gateway: **---gateway ip---**

Name Servers: 8.8.8.8

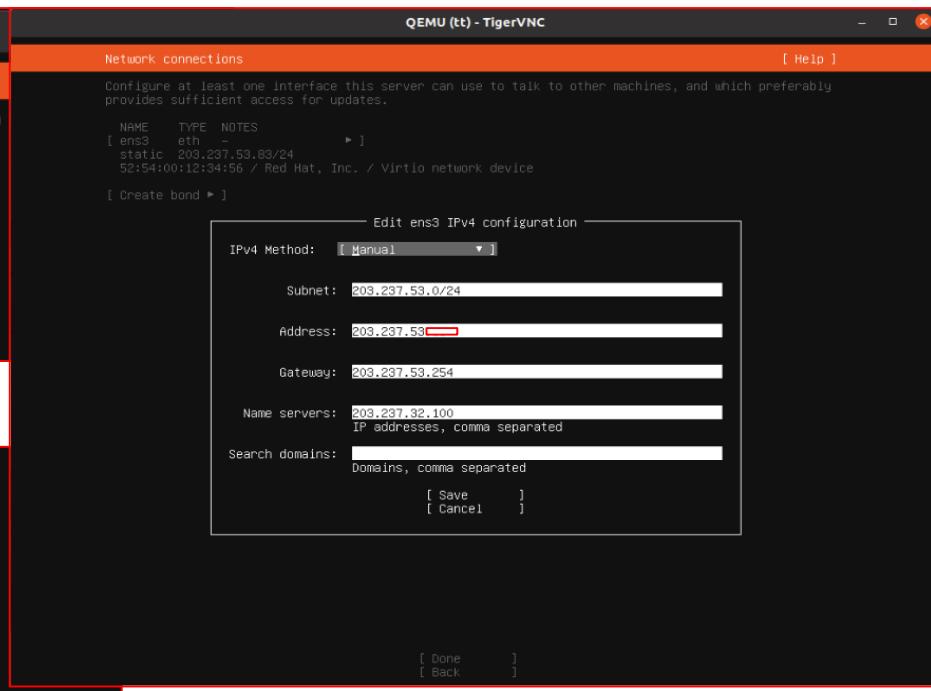
1. Select network device 'Edit IPv4'



2. Select 'Manual'



3. Enter the network info.





#3 - NUC: Making VM with KVM

- Installation Completed (control with 'Enter key' and 'Arrow keys')

When 'installation completed' message is shown, terminate the VM

```
$ sudo killall -9 qemu-system-x86_64
```

boot VM again (mac should be different from others).

```
$sudo kvm -m [memory capacity] -name [name] -smp cpus=[#cpu],maxcpus= [#maxcpu] -device virtio-net-pci,netdev=net0 -netdev tap,id=net0,ifname= [tap_name],script=no -boot d [name].img -vnc :[#] -daemonize
```

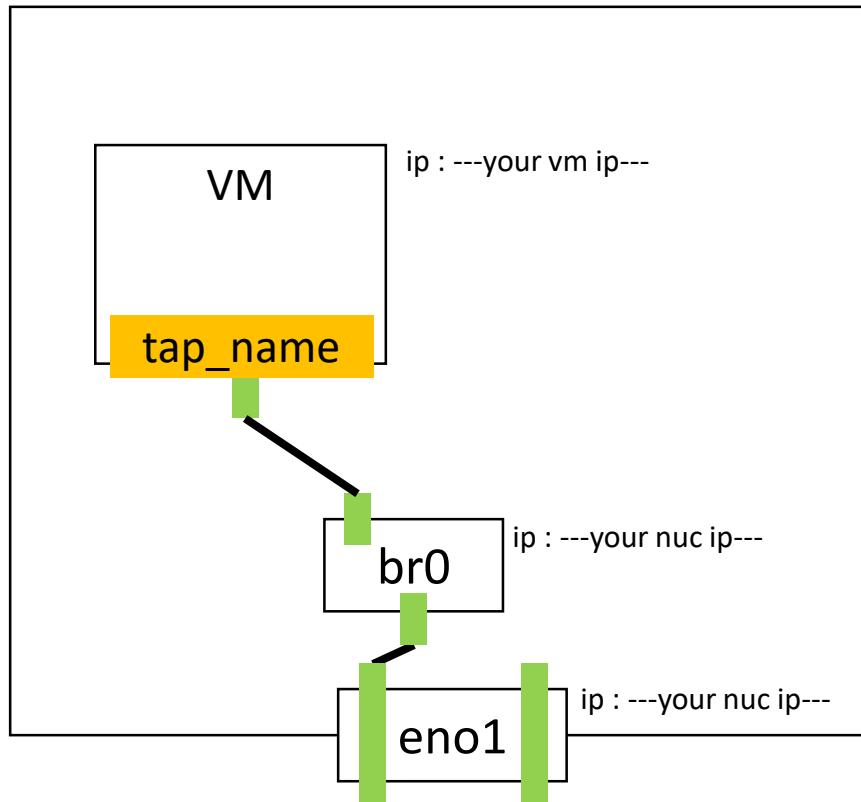
```
$ sudo kvm -m 1024 -name tt -smp cpus=2,maxcpus=2 -device virtio-net-pci,netdev=net0 -netdev tap,id=net0,ifname=vport_vFunction,script=no -boot d vFunction20.img -vnc :5 -daemonize
```

#5 - OVS connects with KVM

- Check situation

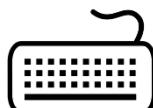


```
root@nuc:~# ovs-vsctl show
3bb93923-3eac-420a-9da9-9143aff14209
    Bridge "br0"
        Port "br0"
            Interface "br0"
                type: internal
        Port vport_vFunction
            Interface vport_vFunction
    ovs_version: "2.0.2"
```



#6 - NUC: Installing ssh in VM

- Don't forget to install ssh in VM



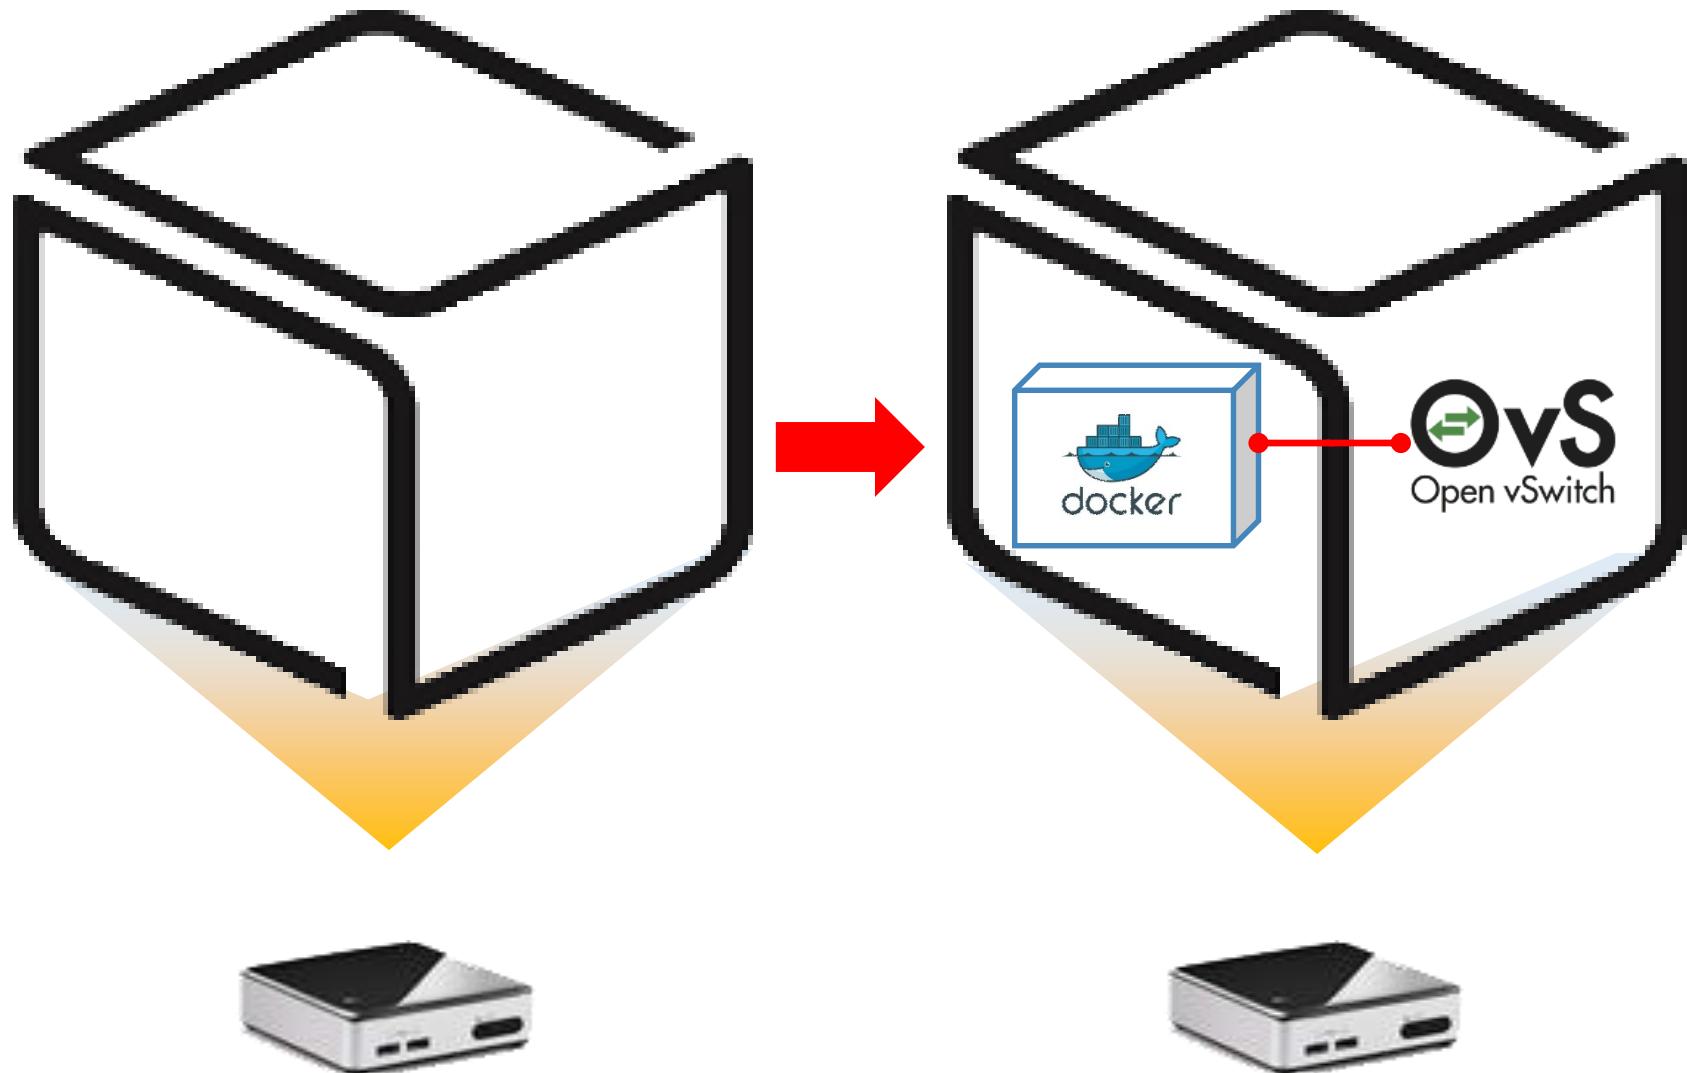
In VMs,

```
$sudo apt-get update  
$sudo apt-get install net-tools ssh -y
```

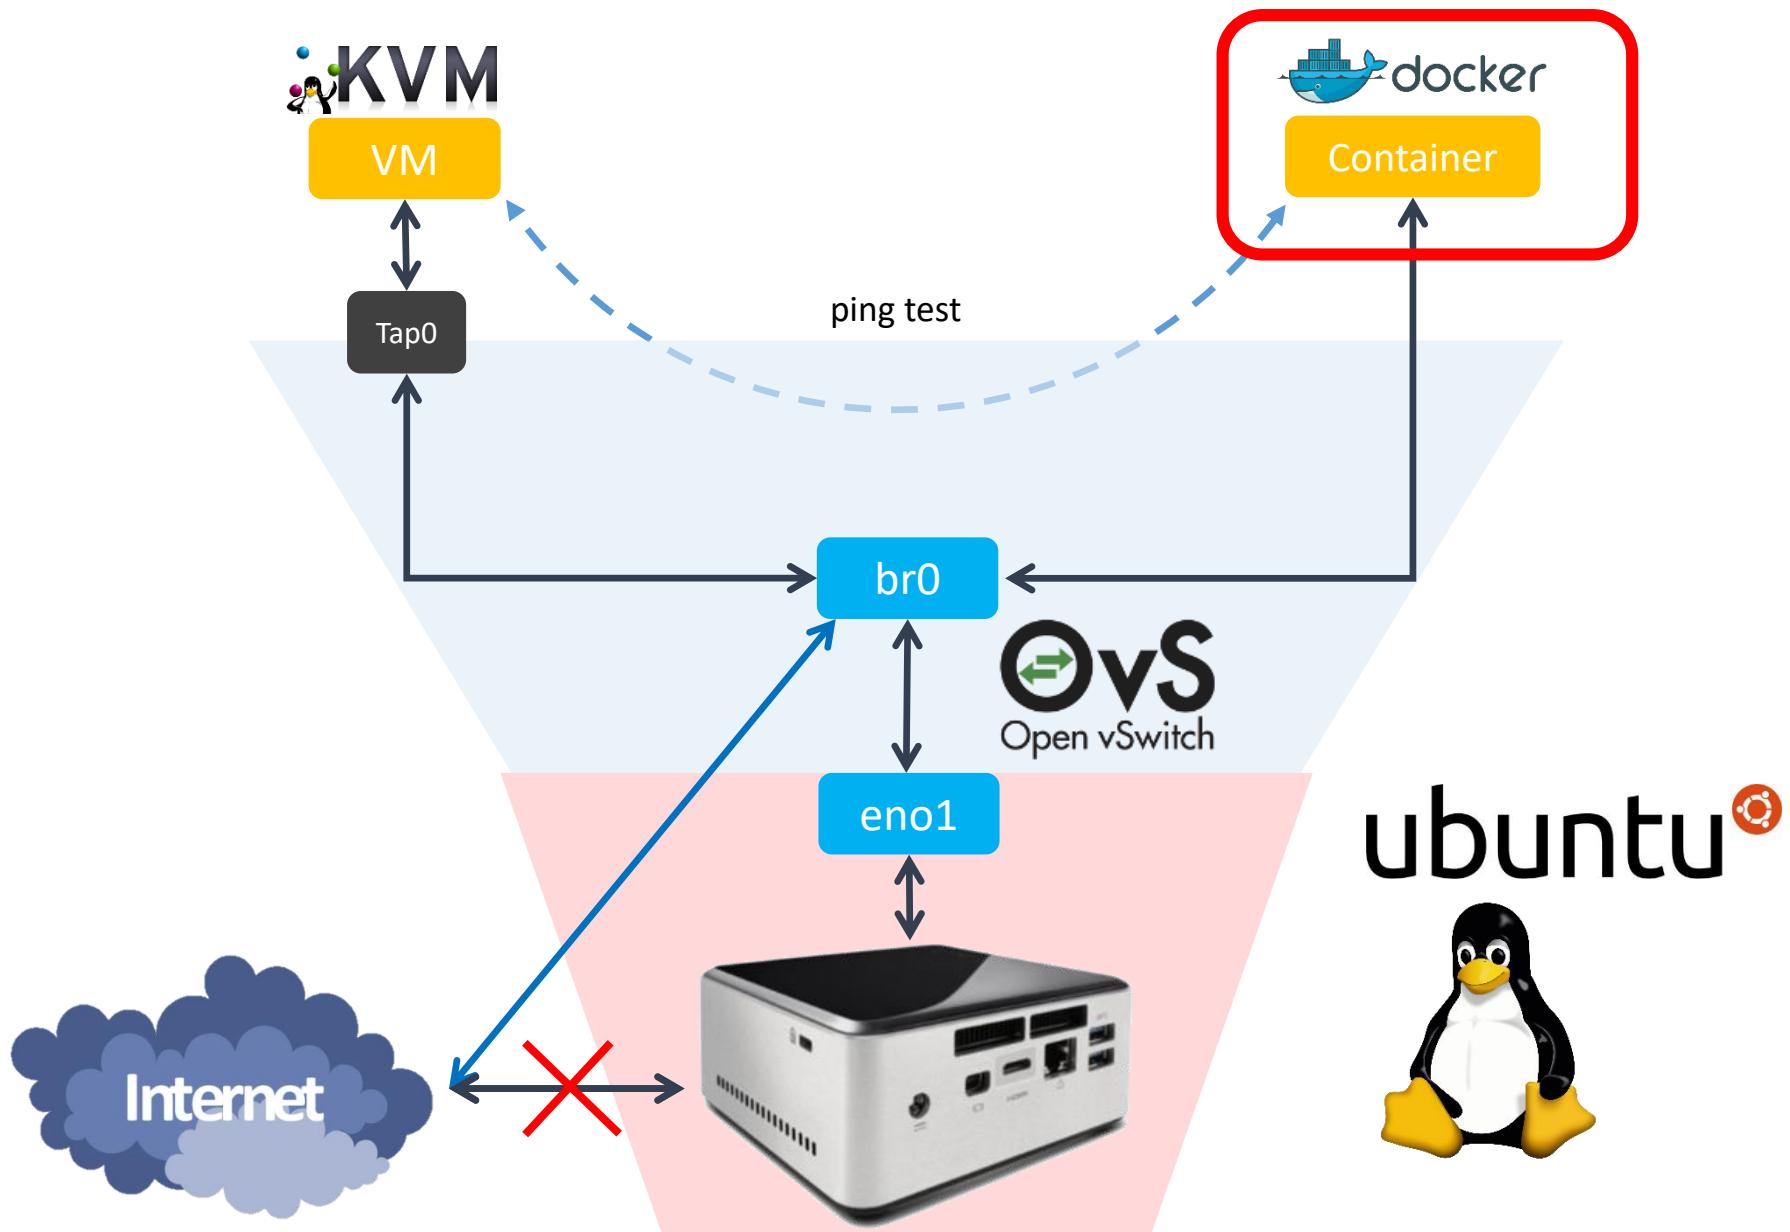
```
nuc@nuc:~$ ssh vbox@192.168.0.3  
The authenticity of host '192.168.0.3 (192.168.0.3)' can't be established.  
ECDSA key fingerprint is da:c5:2c:53:5a:6f:b4:3c:03:02:04:f3:6a:17:ca:ab.  
Are you sure you want to continue connecting (yes/no)? yes  
Warning: Permanently added '192.168.0.3' (ECDSA) to the list of known hosts.  
vbox@192.168.0.3's password: █
```

Docker Container connected via OVS

- Goal of this section



#7 - Making a Docker Container



#7 - Making a Docker Container

- Docker installation

In NUC,

Docker installation.



// Set up the repository:

// Install packages to allow apt to use a repository over HTTPS

```
$ sudo apt-get update && sudo apt-get install -y apt-transport-https ca-certificates curl software-properties-common  
gnupg2
```

// Add Docker's official GPG key:

```
$ curl -fsSL https://download.docker.com/linux/ubuntu/gpg | sudo apt-key --keyring /etc/apt/trusted.gpg.d/docker.gpg  
add -
```

// Add the Docker apt repository:

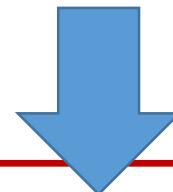
```
$ sudo add-apt-repository \  
$ "deb [arch=amd64] https://download.docker.com/linux/ubuntu \  
$ $(lsb_release -cs) \  
$ stable"
```

// Install Docker CE

```
$ sudo apt-get update && sudo apt-get install -y \  
$ containerd.io=1.2.13-2 \  
$ docker-ce=5:19.03.11~3-0~ubuntu-$(lsb_release -cs) \  
$ docker-ce-cli=5:19.03.11~3-0~ubuntu-$(lsb_release -cs)
```

// Create /etc/docker

```
$ sudo mkdir -p /etc/docker
```



cont'd on next page

#7 - Making a Docker Container

- Docker installation

In NUC,

Docker installation (cont'd)



```
//continued  
// Set up the Docker daemon  
$ cat <<EOF | sudo tee /etc/docker/daemon.json  
$ {  
$     "exec-opts": ["native.cgroupdriver=systemd"],  
$     "log-driver": "json-file",  
$     "log-opt": {  
$         "max-size": "100m"  
$     },  
$     "storage-driver": "overlay2"  
$ }  
$ EOF
```

```
// Create /etc/systemd/system/docker.service.d  
$ sudo mkdir -p /etc/systemd/system/docker.service.d  
  
$ sudo systemctl daemon-reload  
$ sudo systemctl enable docker  
$ sudo systemctl restart docker  
$ sudo systemctl restart docker.socket
```

#7 - Making a Docker Container

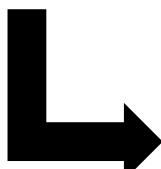
- Docker installation

In NUC,



Now Installation completed. Let's Check it

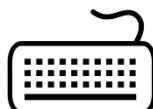
```
$ sudo docker run hello-world
```



```
Hello from Docker.  
This message shows that your installation appears to be working correctly.  
  
To generate this message, Docker took the following steps:  
1. The Docker client contacted the Docker daemon.  
2. The Docker daemon pulled the "hello-world" image from the Docker Hub.  
3. The Docker daemon created a new container from that image which runs the  
executable that produces the output you are currently reading.  
4. The Docker daemon streamed that output to the Docker client, which sent it  
to your terminal.  
  
To try something more ambitious, you can run an Ubuntu container with:  
$ docker run -it ubuntu bash  
  
Share images, automate workflows, and more with a free Docker Hub account:  
https://hub.docker.com  
  
For more examples and ideas, visit:  
https://docs.docker.com/userguide/
```

#7 - Making a Docker Container

- Make container



Run docker container. (**Remember the container_name)

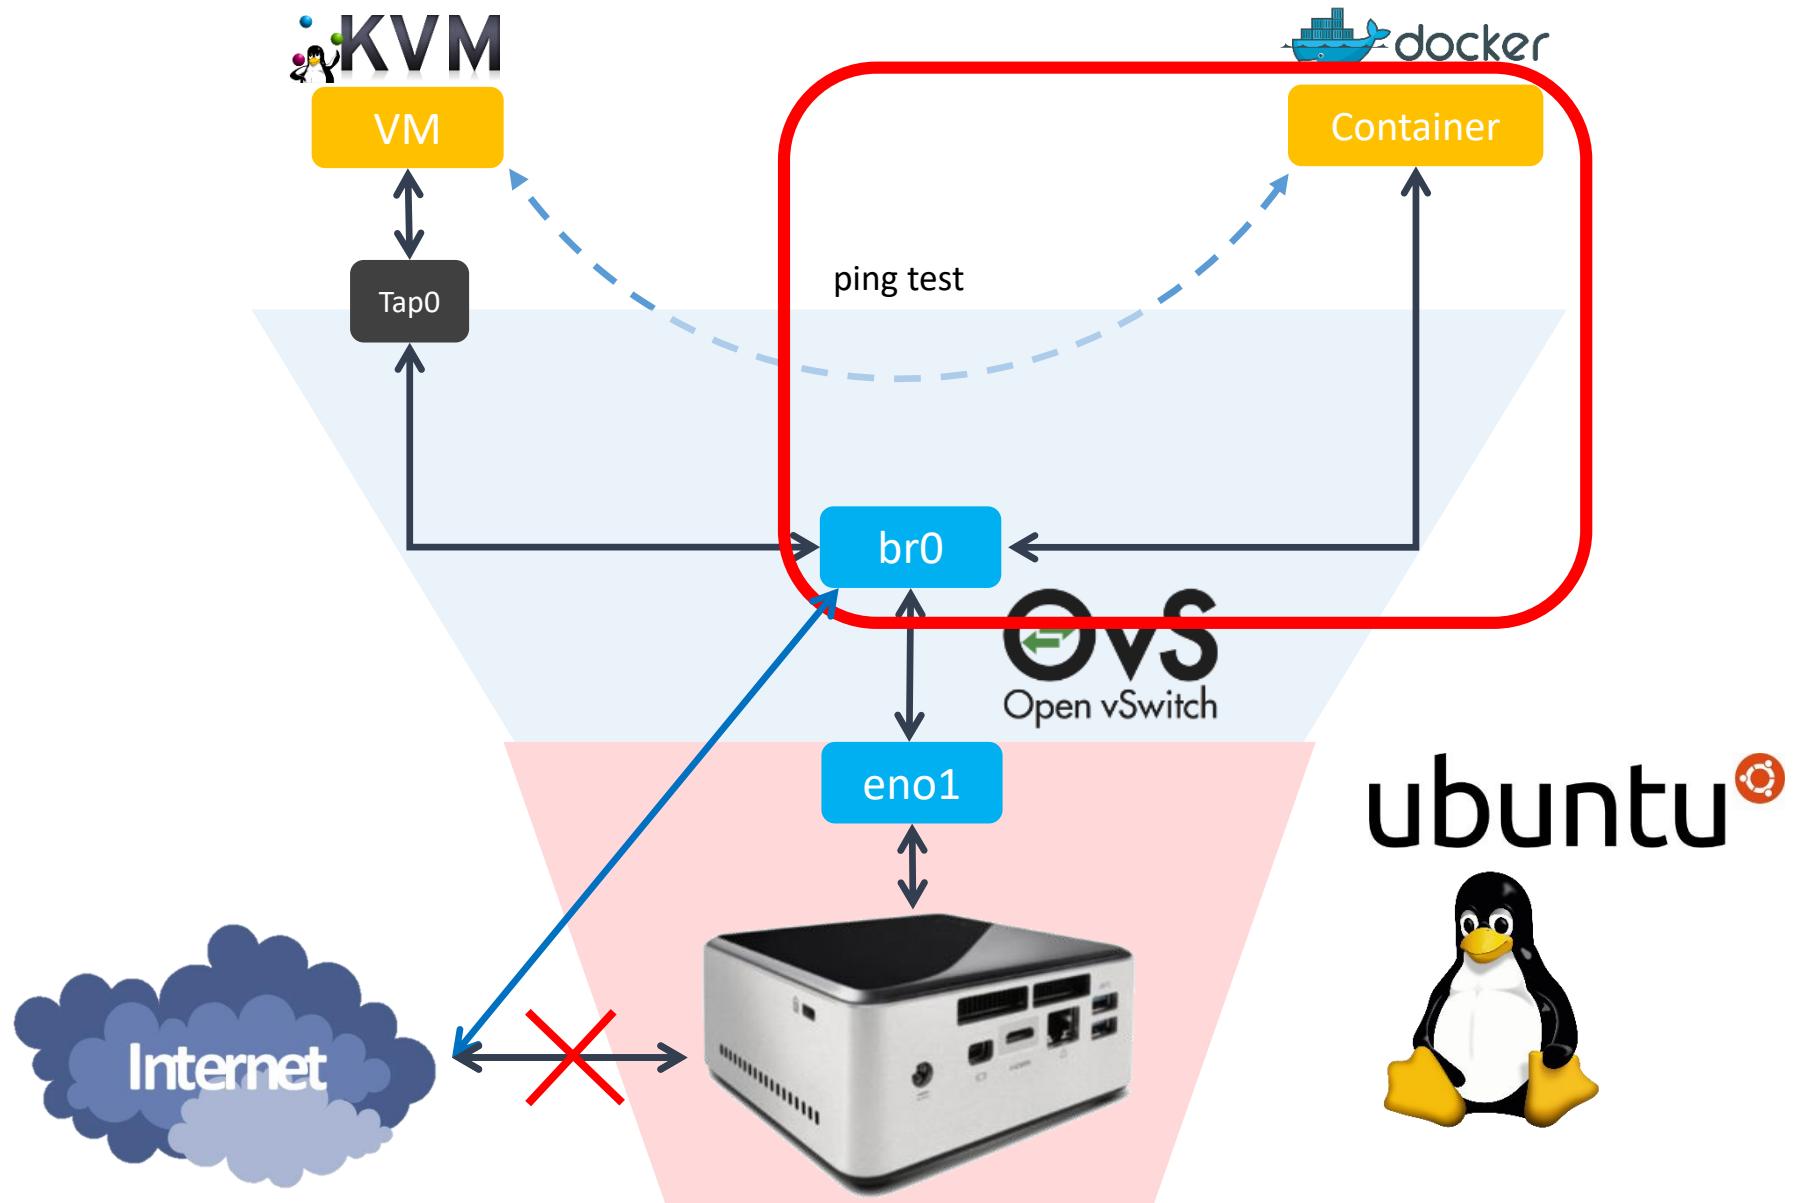
```
$sudo docker run -it --net=none --name [container_name] ubuntu /bin/bash
```

```
nuc@nuc:~$ docker run -it --net=none --name c1 ubuntu /bin/bash
root@8346684676d8:/#
```

- ⌘ctrl + p, q → detach docker container
- ⌘docker attach [container_name] → get into docker container console

#8 - Connect Docker Container

- Connect with OVS bridge



#8 - Connect docker Container

- Connect with OVS bridge



Install OVS-docker utility in host machine. (Not in inside of Docker container.)

```
$sudo ovs-docker add-port br0 veno1 [container_name] --ipaddress=---your docker ip---/24 --gateway=---gateway ip---
```

```
$sudo docker attach [container_name] //enter to docker container  
#apt-get update  
#apt-get install net-tools  
#apt-get install iputils-ping
```

#9 – Keep Docker network configuration

- /etc/rc.local



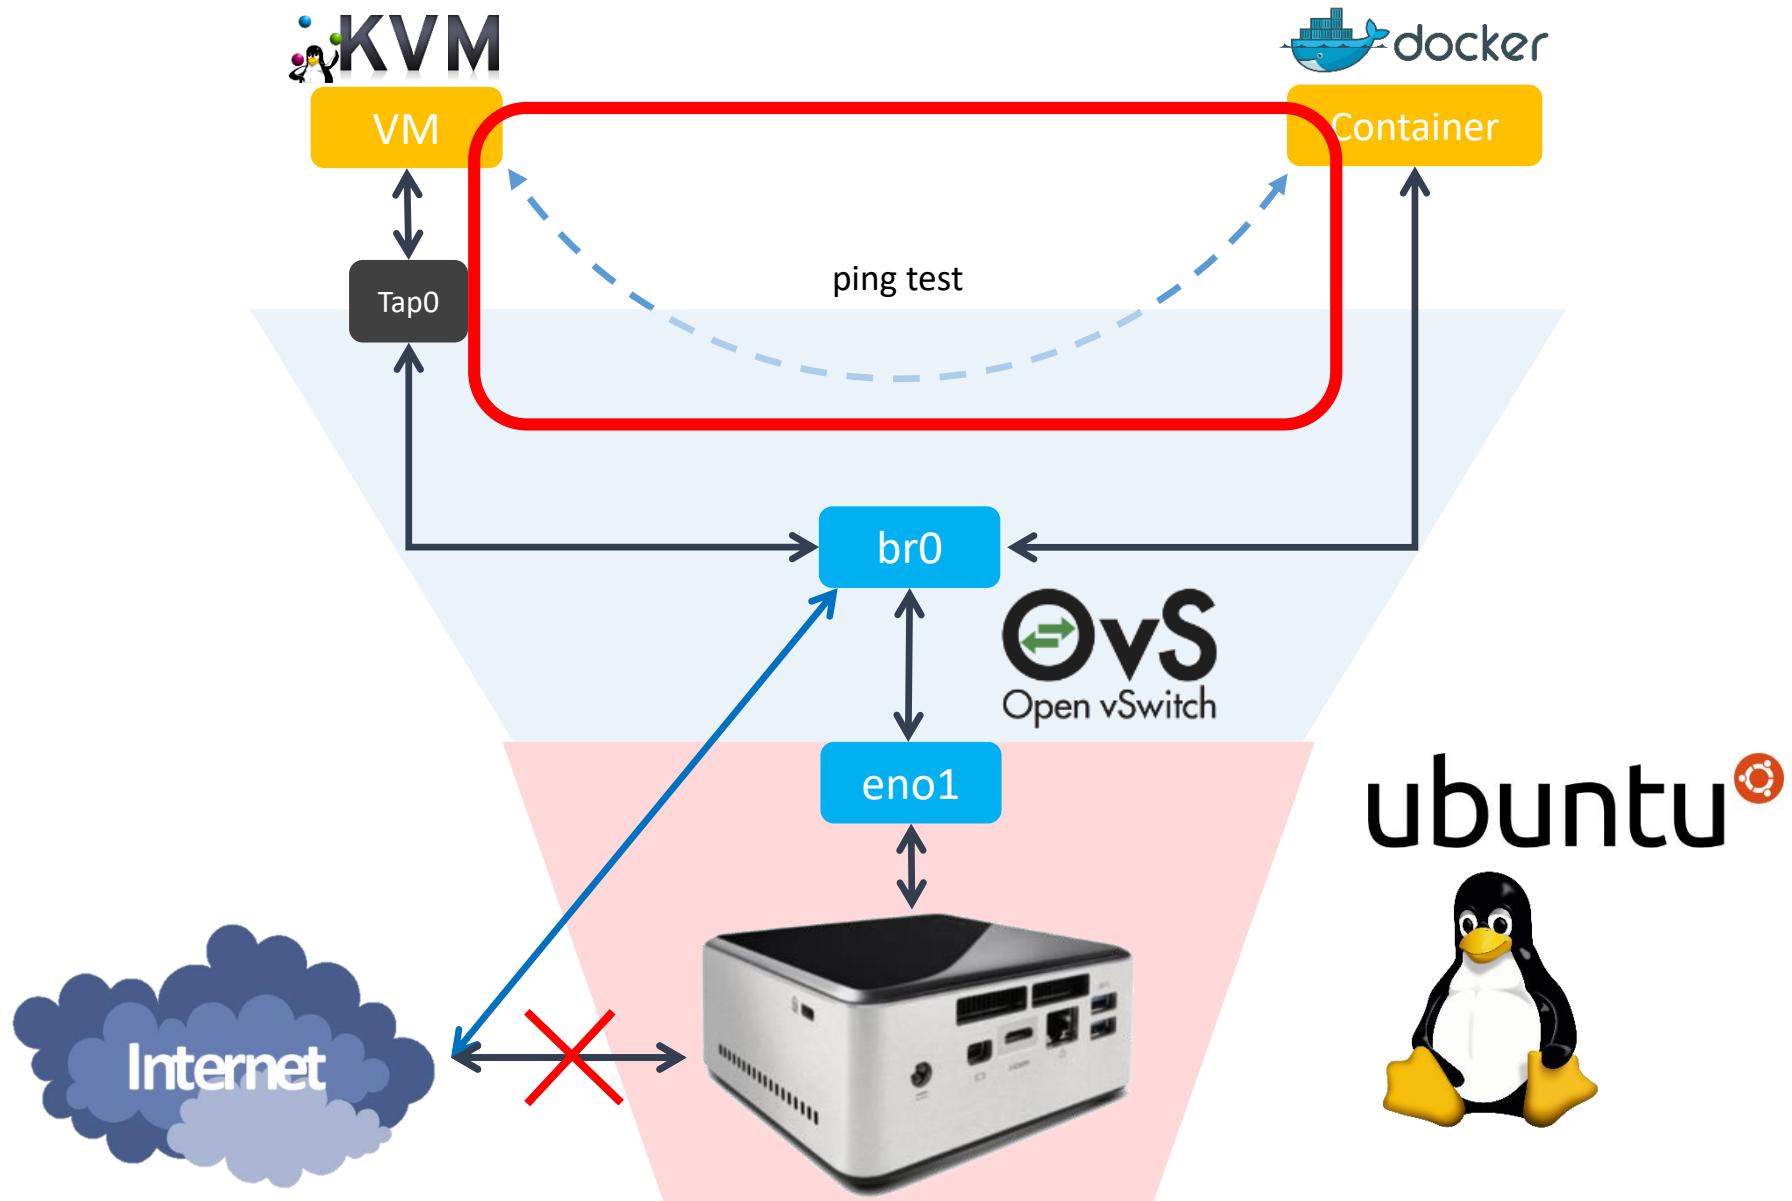
Modify /etc/rc.local

```
$sudo vi /etc/rc.local
```

```
#!/bin/bash
docker start [container_name]
ovs-docker del-port br0 veno1 [containerName]
ovs-docker add-port br0 veno1 [container_name] --ipaddress=---your docker ip---/24 --gateway=---gateway ip---
```

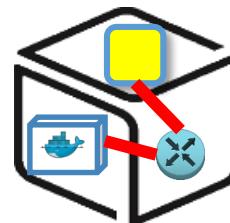
Whenever NUC is rebooted,
network configuration of Docker container is initialized
by executing commands in **rc.local**

#10 – Check connectivity: VM & Container



#10 - Check connectivity: VM & Container

-Check connectivity with ping command



```
root@nuc:/usr/bin# ovs-docker add-port br0 eth0 docker1 --ipaddress=210.125.      /24 --gateway=210.125
root@nuc:/usr/bin# docker attach docker1

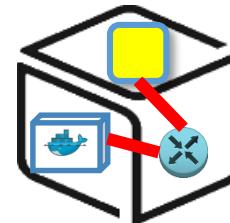
root@b8c3bab8204b:/# ifconfig
eth0      Link encap:Ethernet  HWaddr ae:e5:9c:cc:88:b7
          inet  addr:210.125         Bcast:0.0.0.0  Mask:255.255.255.0
          inet6 addr: fe80::ace5:9cff:fecc:88b7/64 Scope:Link
            UP BROADCAST RUNNING MULTICAST  MTU:1500 Metric:1
            RX packets:120 errors:0 dropped:0 overruns:0 frame:0
            TX packets:8 errors:0 dropped:0 overruns:0 carrier:0
            collisions:0 txqueuelen:1000
            RX bytes:8842 (8.8 KB)   TX bytes:648 (648.0 B)

lo       Link encap:Local Loopback
          inet  addr:127.0.0.1  Mask:255.0.0.0
          inet6 addr: ::1/128 Scope:Host
            UP LOOPBACK RUNNING  MTU:65536 Metric:1
            RX packets:0 errors:0 dropped:0 overruns:0 frame:0
            TX packets:0 errors:0 dropped:0 overruns:0 carrier:0
            collisions:0 txqueuelen:0
            RX bytes:0 (0.0 B)   TX bytes:0 (0.0 B)

root@b8c3bab8204b:/# ping google.com
PING google.com (216.58.221.238) 56(84) bytes of data.
64 bytes from hkg07s21-in-f14.1e100.net (216.58.221.238): icmp_seq=1 ttl=52 time=41.3 ms
64 bytes from hkg07s21-in-f14.1e100.net (216.58.221.238): icmp_seq=2 ttl=52 time=41.3 ms
^C
--- google.com ping statistics ---
2 packets transmitted, 2 received, 0% packet loss, time 1001ms
rtt min/avg/max/mdev = 41.306/41.343/41.380/0.037 ms
```

#10 - Check connectivity: VM & Container

-Check connectivity with ping command



```

root@b8c3bab8204b:/# ifconfig
eth0      Link encap:Ethernet HWaddr a2:86:d9:c2:33
          inet addr:192.168.0.0.0 Mask:255.255.255.0
          inet6 addr: fe80::a086:d9ff:fecc:337b/64 Brd:192.168.0.255 Scope:Link
          UP BROADCAST RUNNING MULTICAST MTU:1500 Metric:1
          RX packets:136 errors:0 dropped:0 overruns:0
          TX packets:13 errors:0 dropped:0 overruns:0
          collisions:0 txqueuelen:1000
          RX bytes:10448 (10.4 KB) TX bytes:1043 (1)

lo        Link encap:Local Loopback
          inet addr:127.0.0.1 Mask:255.0.0.0
          inet6 addr: ::1/128 Scope:Host
          UP LOOPBACK RUNNING MTU:65536 Metric:1
          RX packets:0 errors:0 dropped:0 overruns:0
          TX packets:0 errors:0 dropped:0 overruns:0
          collisions:0 txqueuelen:0
          RX bytes:0 (0.0 B) TX bytes:0 (0.0 B)

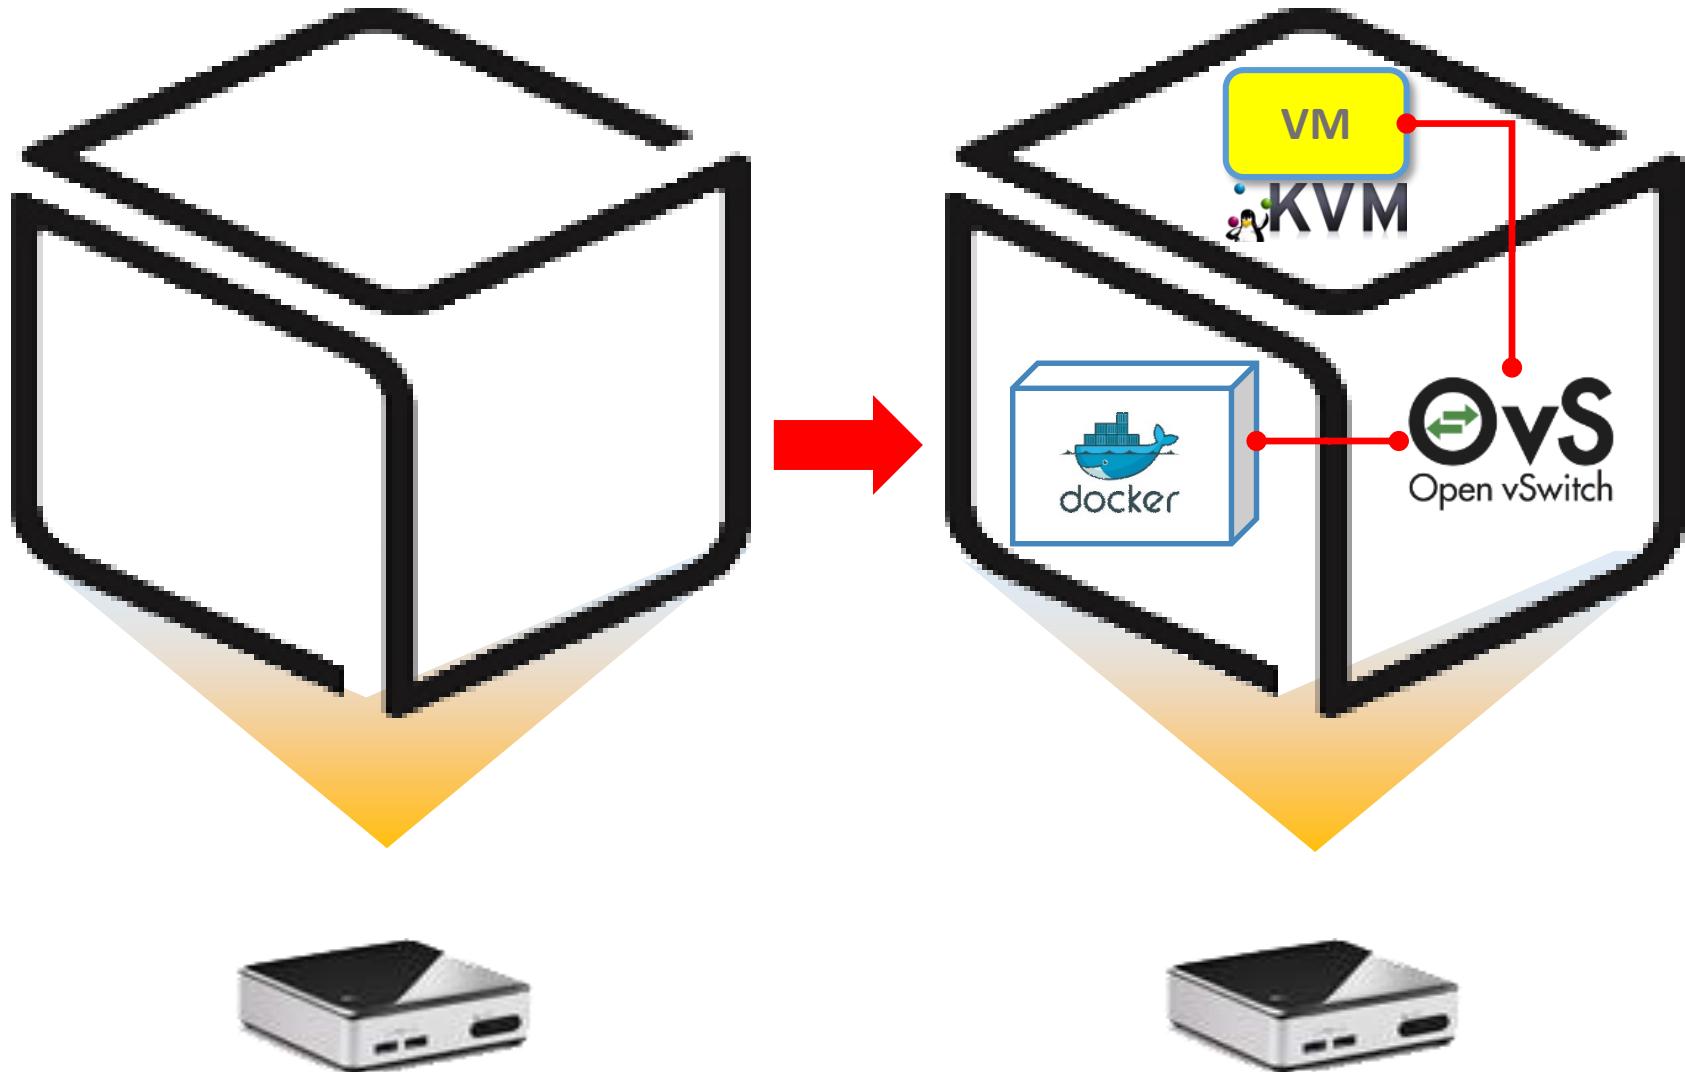
root@b8c3bab8204b:/# ping google.com
PING google.com (216.58.221.238) 56(84) bytes of data
64 bytes from hkg07s21-in-f238.1e100.net (216.58.221.238): icmp_seq=1 ttl=64 time=1.376 ms
64 bytes from hkg07s21-in-f238.1e100.net (216.58.221.238): icmp_seq=2 ttl=64 time=1.380 ms
^C
--- google.com ping statistics ---
2 packets transmitted, 2 received, 0% packet loss, time 1ms
rtt min/avg/max/mdev = 1.376/1.380/1.384/0.004 ms
root@b8c3bab8204b:/# ping 192.168.0.2
PING 192.168.0.2 (192.168.0.2) 56(84) bytes of data.
64 bytes from 192.168.0.2: icmp_seq=1 ttl=64 time=1.376 ms
64 bytes from 192.168.0.2: icmp_seq=2 ttl=64 time=1.380 ms
64 bytes from 192.168.0.2: icmp_seq=3 ttl=64 time=1.384 ms
^C
--- 192.168.0.2 ping statistics ---
3 packets transmitted, 3 received, 0% packet loss, time 2ms
rtt min/avg/max/mdev = 0.651/1.028/1.519/0.365 ms
root@b8c3bab8204b:/# 
```

Docker container

KVM VM

Box Lab: Final Goal (Recap)

Lab #1: Box 47



Lab Review

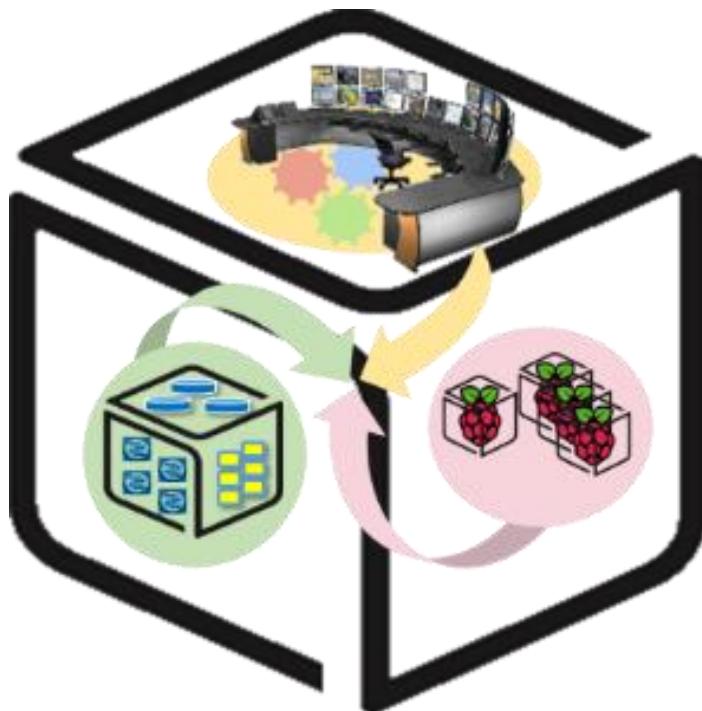


Lab Summary

With Box Lab, you have experimented

1. How to install and configure **Linux OS** into Box (i.e., computer).
2. How to install and configure **OVS (Open vSwitch) virtual switch** inside a Linux Box and configure it.
3. How to create **VMs and Docker containers** inside a Linux Box and then **inter-connect** each of them together and to the Internet.

Thank You for Your Attention Any Questions?



mini@smartx.kr