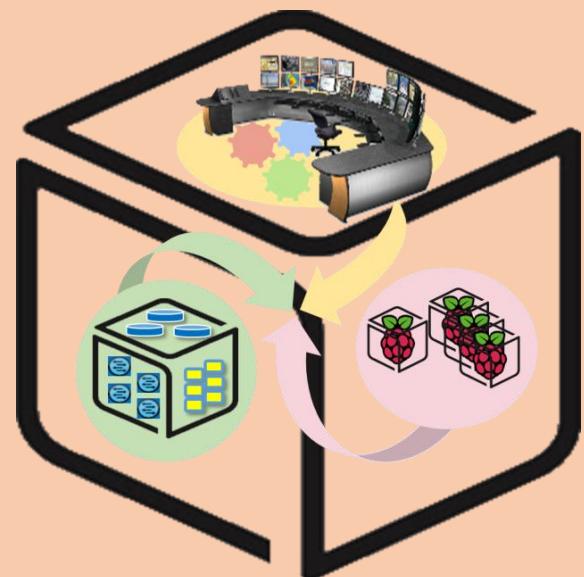


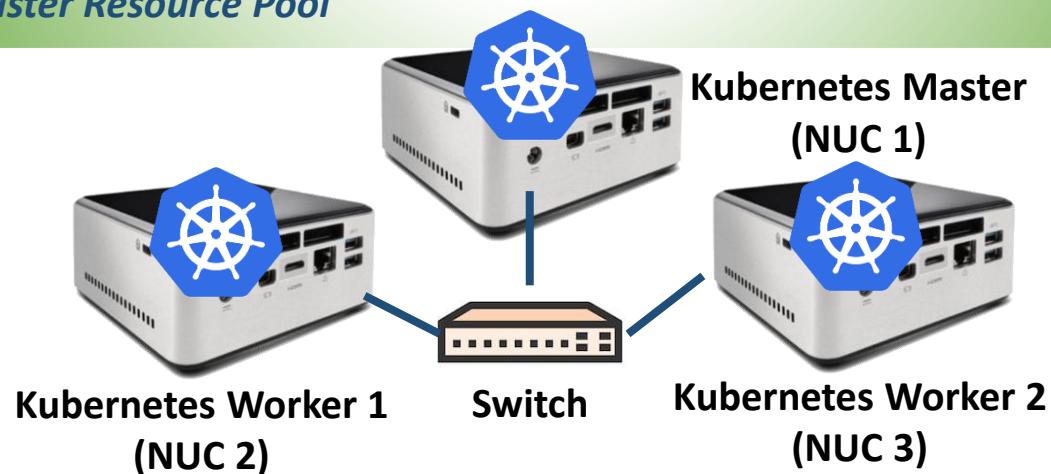
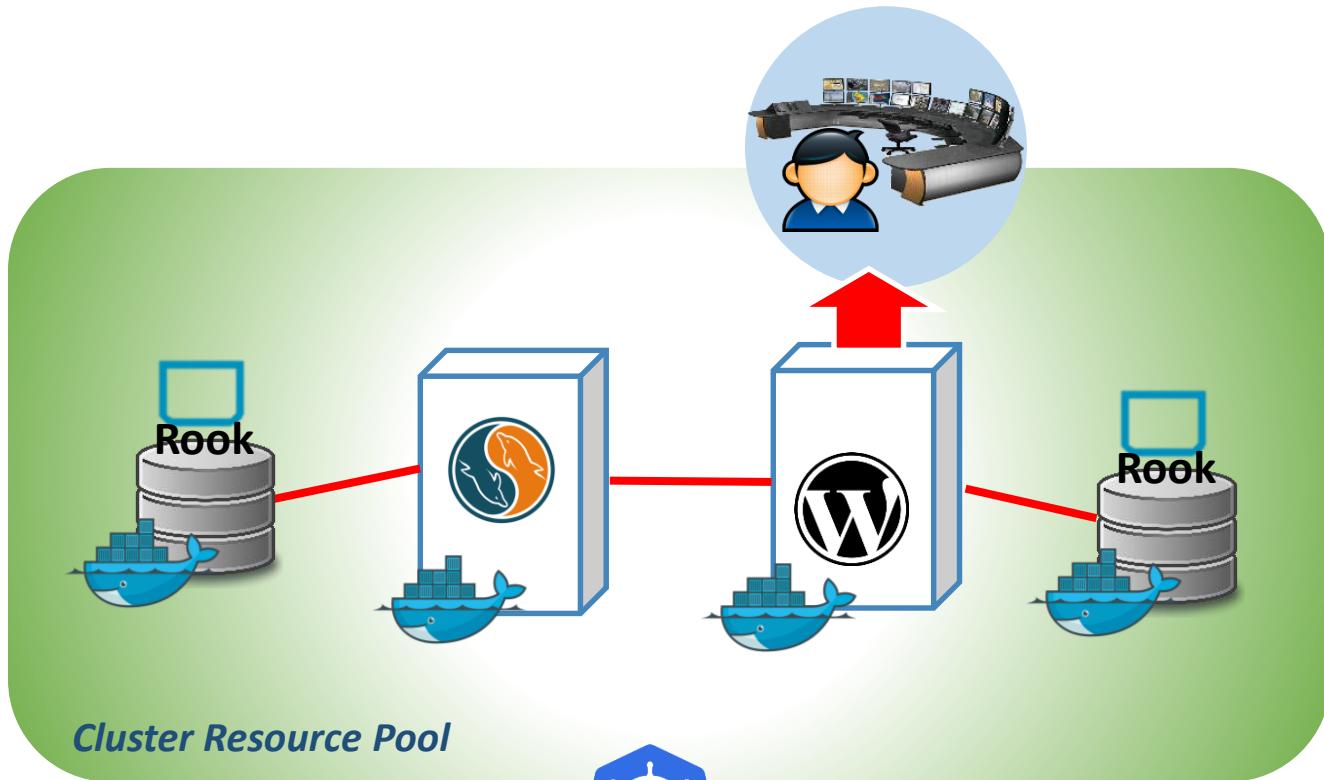
Computer Systems For AI-inspired Cloud Theory & Lab.

Lab #5: Cluster

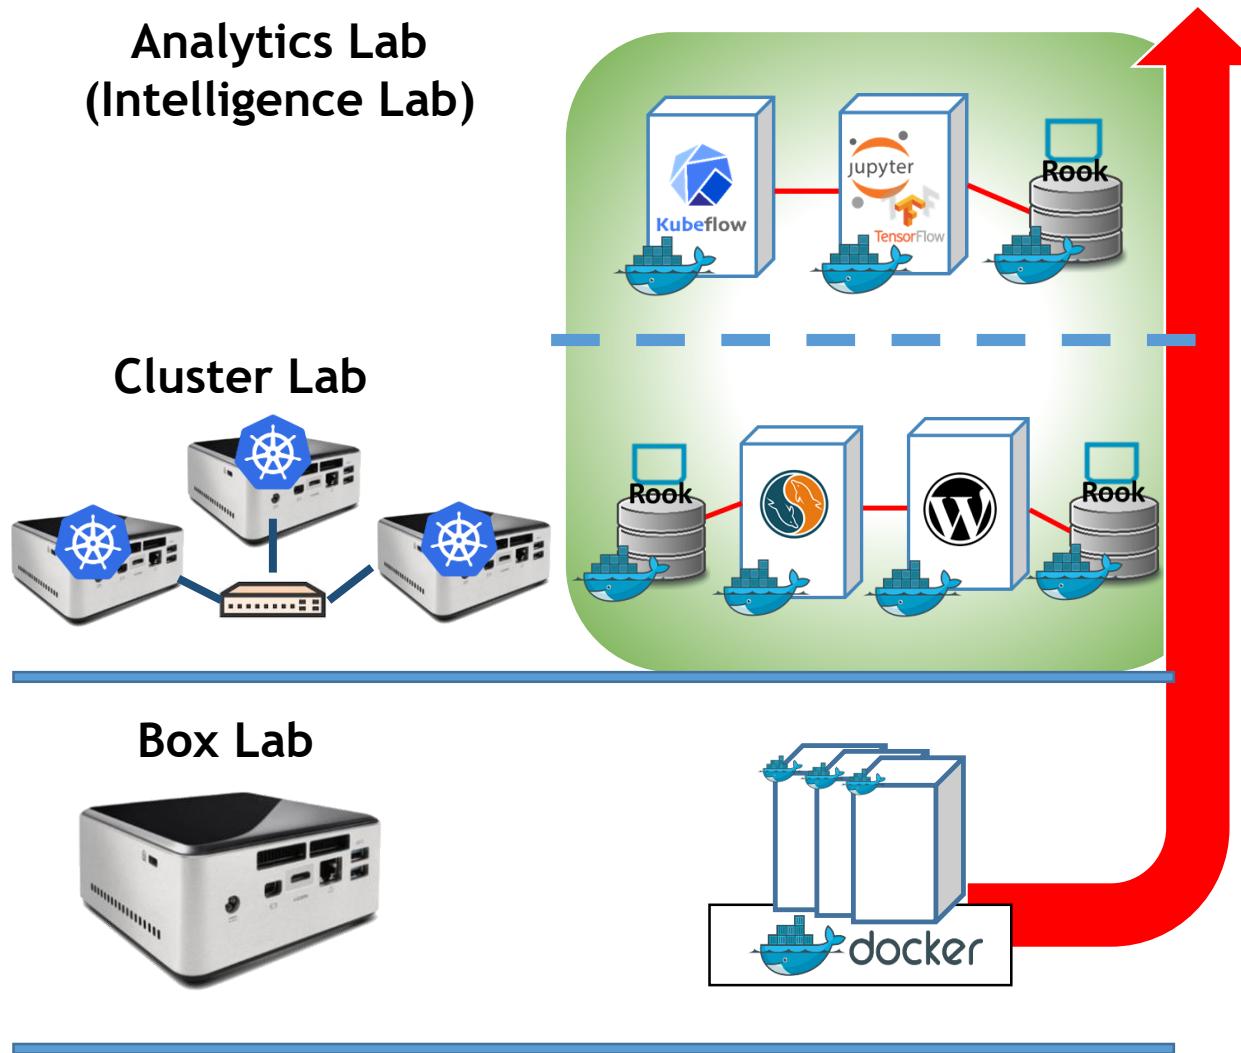


<https://github.com/SmartX-Labs/SmartX-Mini-MOOC>

Cluster Lab: Concept



SmartX Labs #1/#5/#6: Relationship

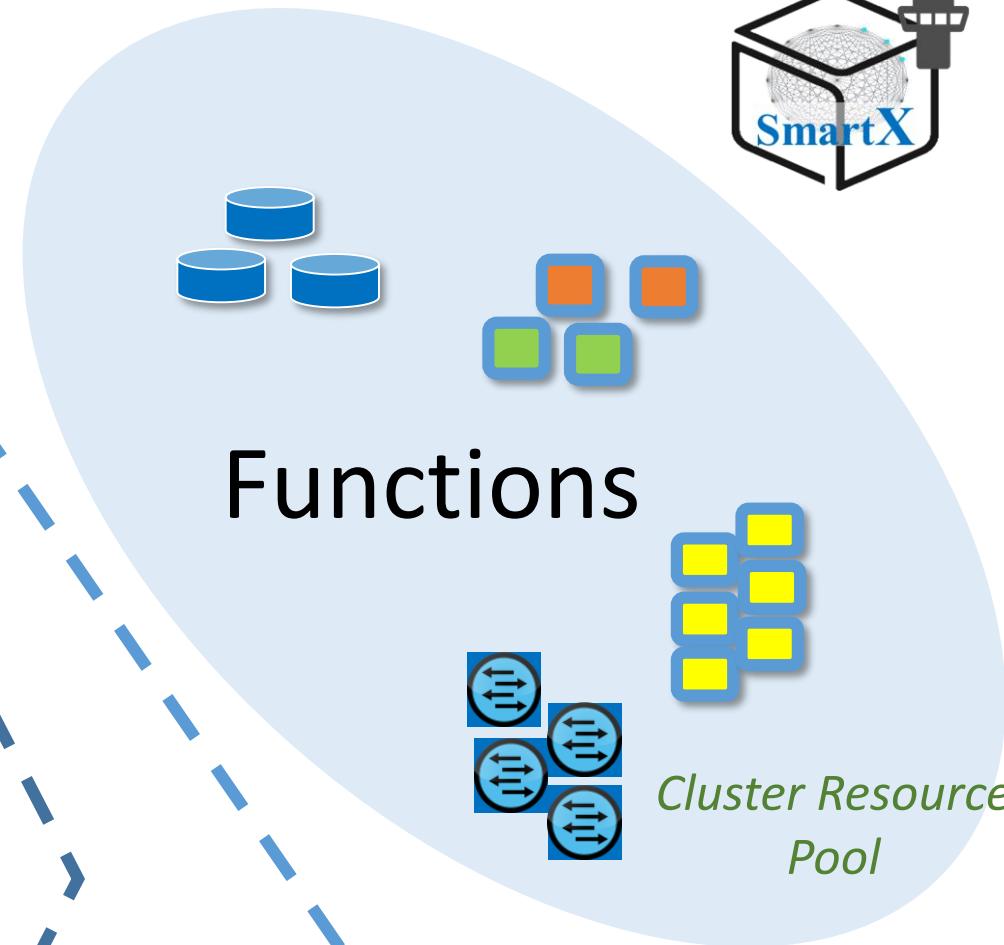
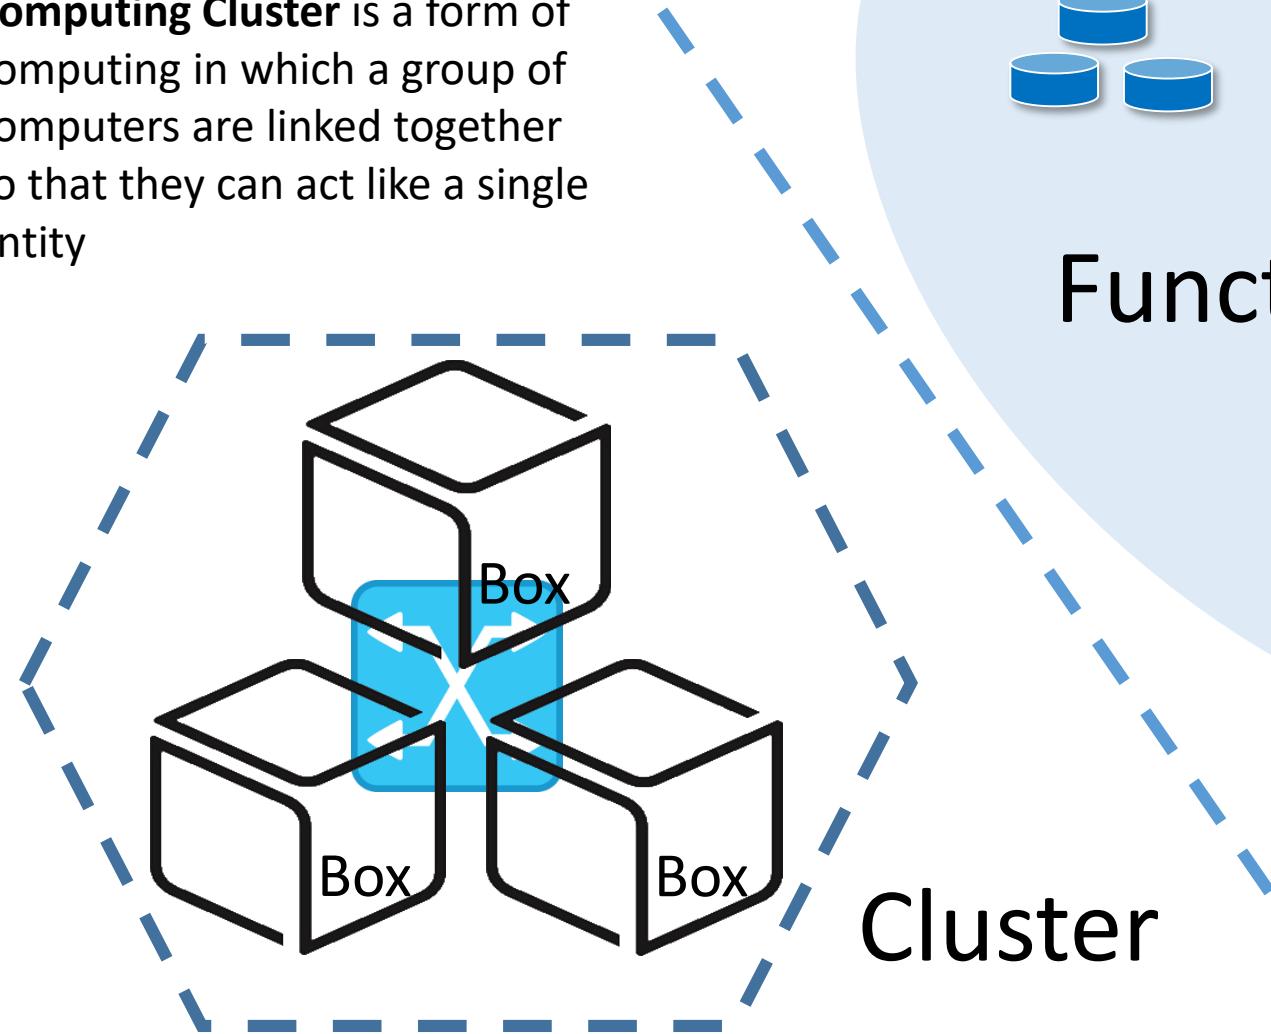


Theory



SmartX Cluster: Inter-connected SmartX Boxes

Computing Cluster is a form of computing in which a group of computers are linked together so that they can act like a single entity

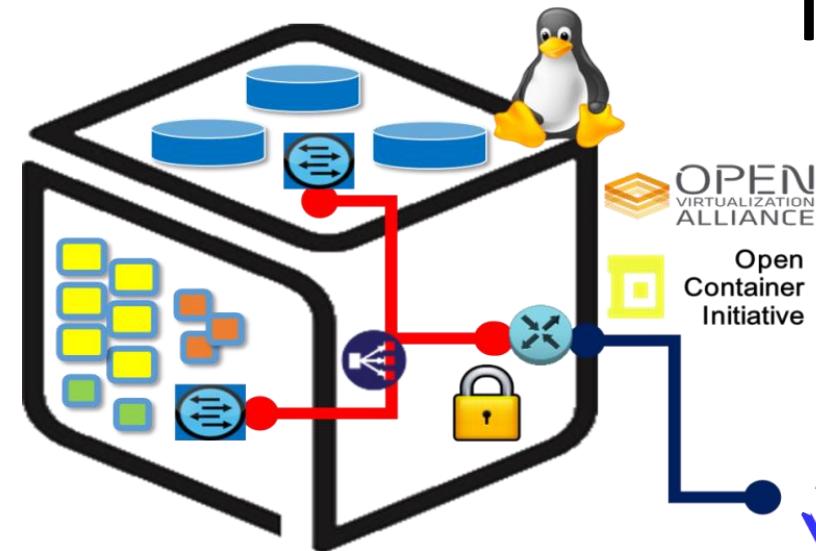




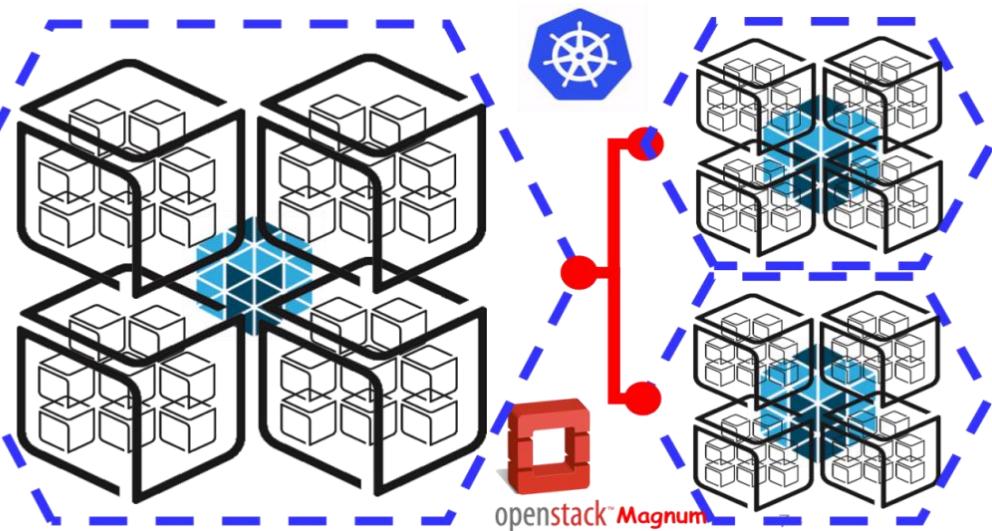
Inter-Connected Functions inside a Box & across Clusters

p+v+c Harmonization Challenge:
p(Baremetal) + v(VM) + c(Container)

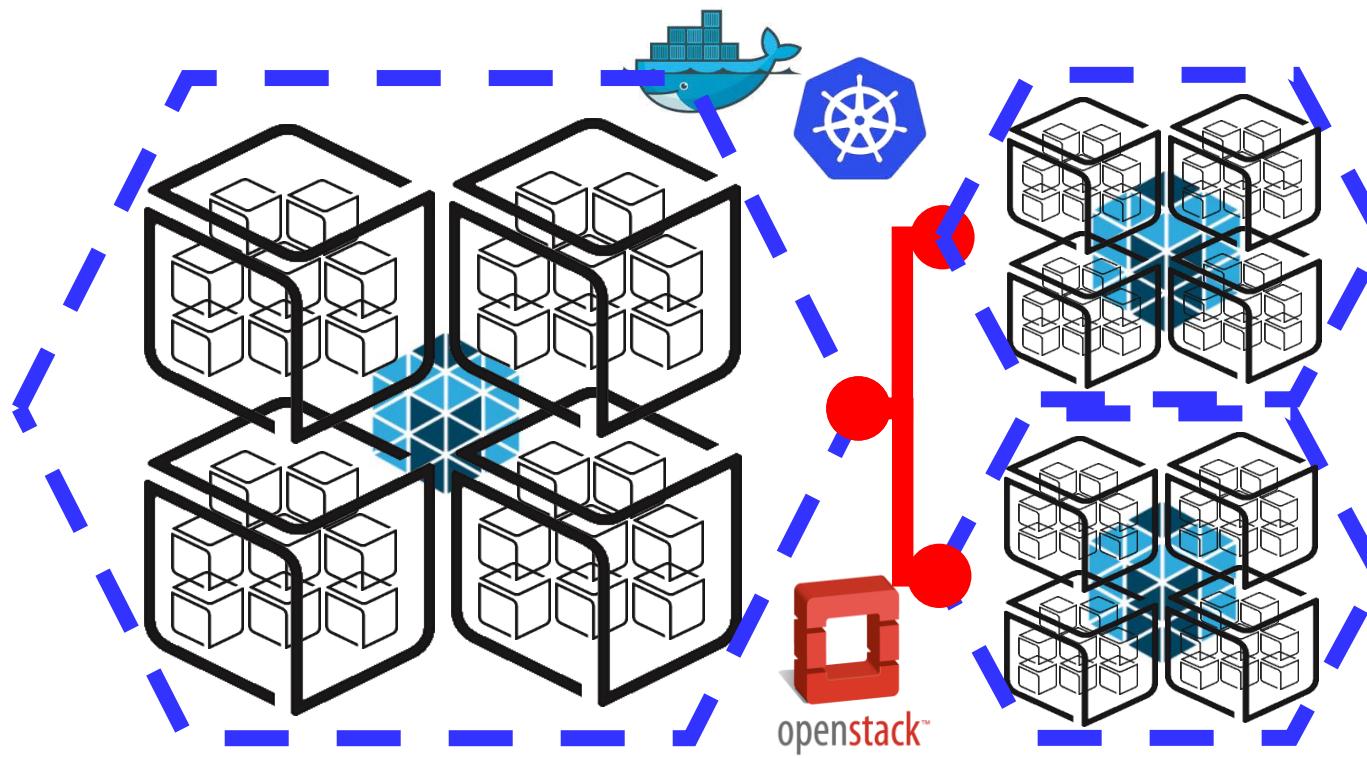
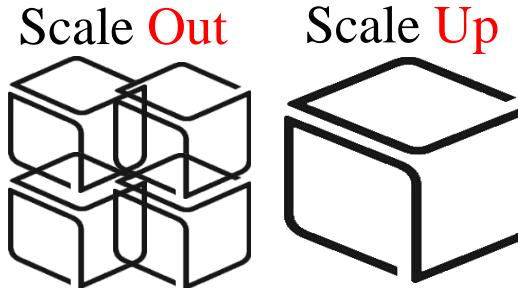
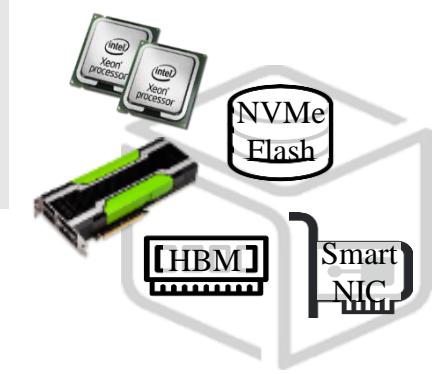
Inside a Box



Across Clusters

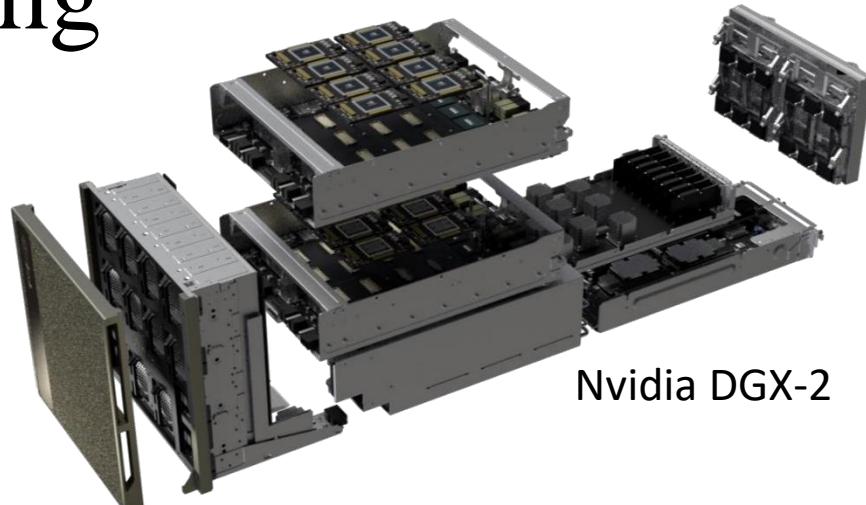


Computer System: Resource Scaling/Pooling with Clustering



Cluster for AI Computing

HPC + HPDA (BigData)
 → AI (ML/DL)



 **Intelligence Center**



GPU Acceleration



GPU Cards

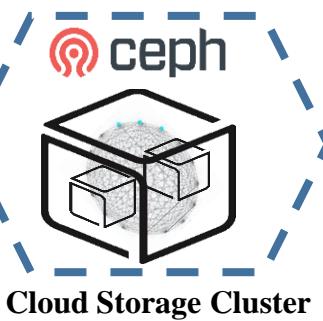
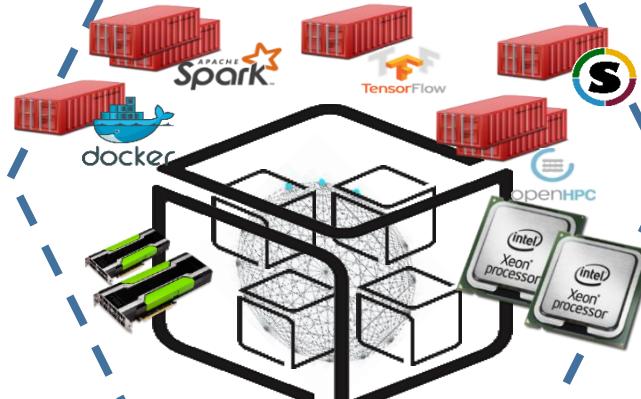
Parallel File System



NVMe SSDs

**Multi-node AI Computing Cluster
 with Optimized DL Tools**

Container Orchestration
 kubernetes



Cloud Storage



Smart NICs

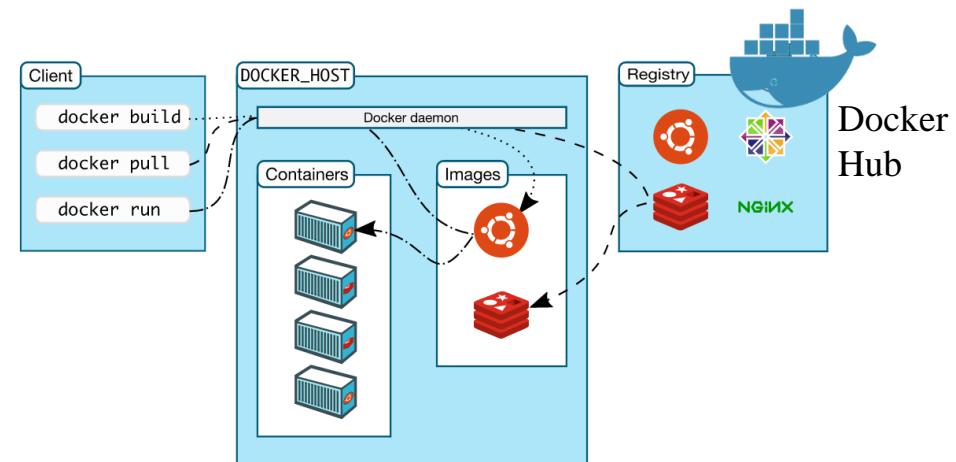
Docker Containers



- **Docker** is an open platform for building, shipping and running distributed applications. It gives programmers, development teams and operations engineers the common toolbox they need to take advantage of the distributed and networked nature of modern applications.



Containers



Since **container uses host kernel**,
OS of host should be **Linux** distribution.

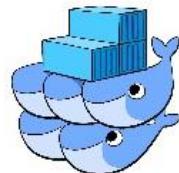
Container Orchestration

Container Orchestration Tools



MESOS

Marathon (Mesosphere)



Docker Swarm



Nomad (HashiCorp)



Kubernetes

Most Popular

Container orchestration refers to the process of organizing the work of individual components and application layers.

Container orchestration engines all allow users to control when containers start and stop, group them into clusters, and coordinate all of the processes that compose an application. Container orchestration tools allow users to guide container deployment and automate updates, health monitoring, and failover procedures.

Container Orchestration: Kubernetes

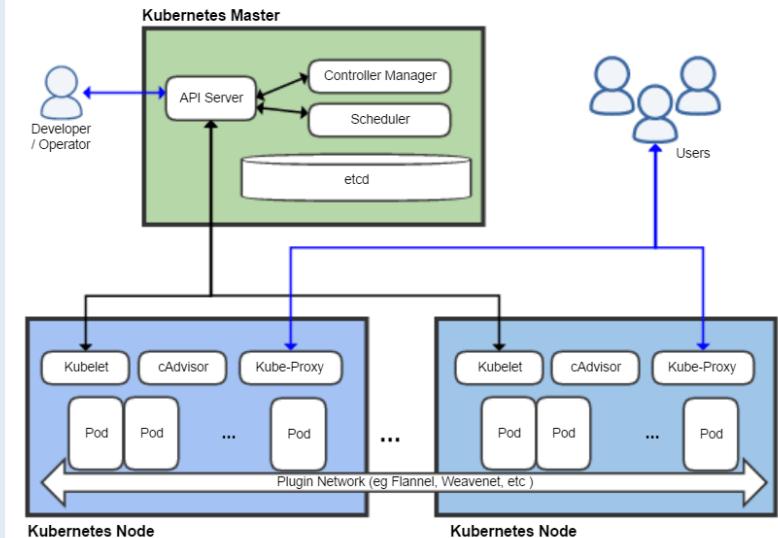
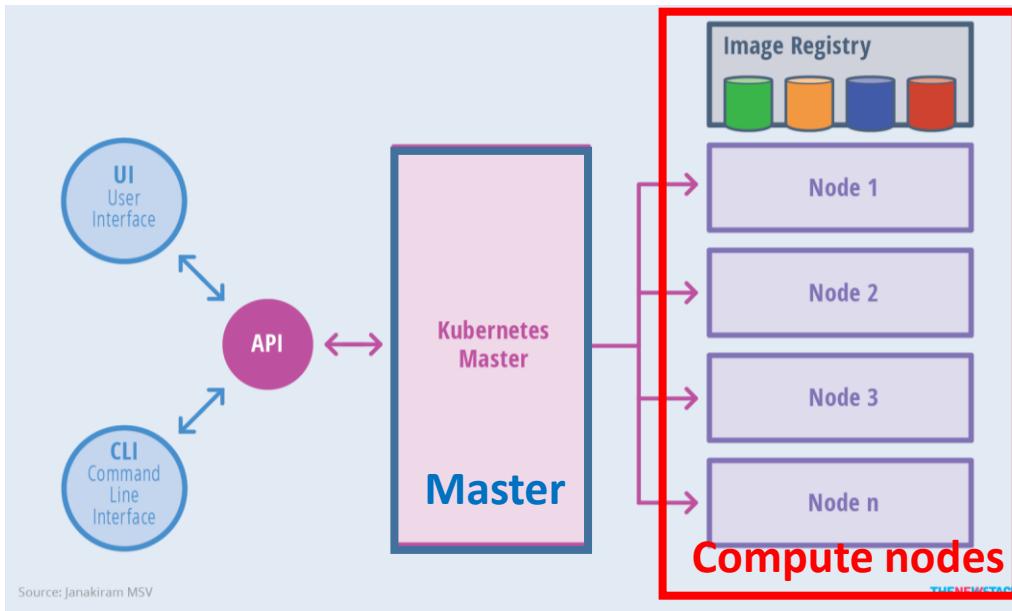
Kubernetes is an open-source system for automating deployment, scaling, and management of containerized applications.



Kubernetes Features

- **Horizontal scaling:** Scale your application up and down with a simple command, with a UI, or automatically based on CPU usage.
- **Self-healing:** Restarts containers that fail, replaces and reschedules containers when nodes die, kills containers that don't respond to your user-defined health check, and doesn't advertise them to clients until they are ready to serve.
- **Service discovery and load balancing:** No need to modify your application to use an unfamiliar service discovery mechanism. Kubernetes gives containers their own IP addresses and a single DNS name for a set of containers, and can load-balance across them.
- **Storage Orchestration:** Automatically mount the storage system of your choice, whether from local storage, a public cloud provider

Container Orchestration: Kubernetes

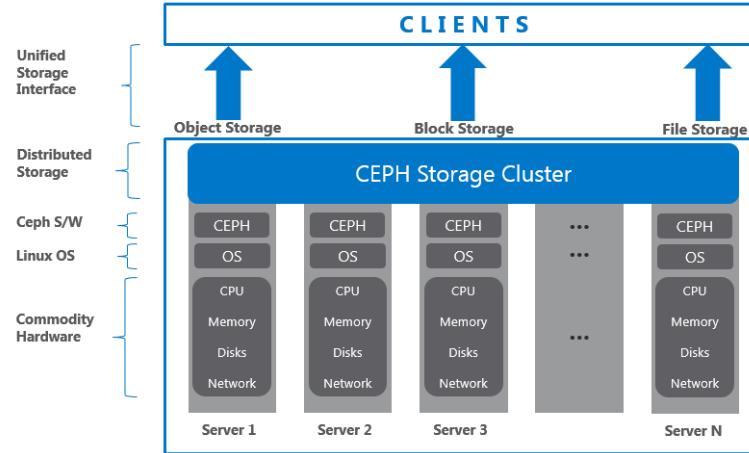


Kubernetes cluster consists of at least one master and multiple **compute nodes**.

The **master** is responsible for exposing the application program interface (**API**), scheduling the deployments and managing the overall cluster.

Pod is consists of one or more containers that are guaranteed to be co-located on the host machine and can share resources. Each pod is assigned a unique IP address within the cluster, which allows applications to use ports without the risk of conflict.

Storage: Ceph and Rook (1/2)

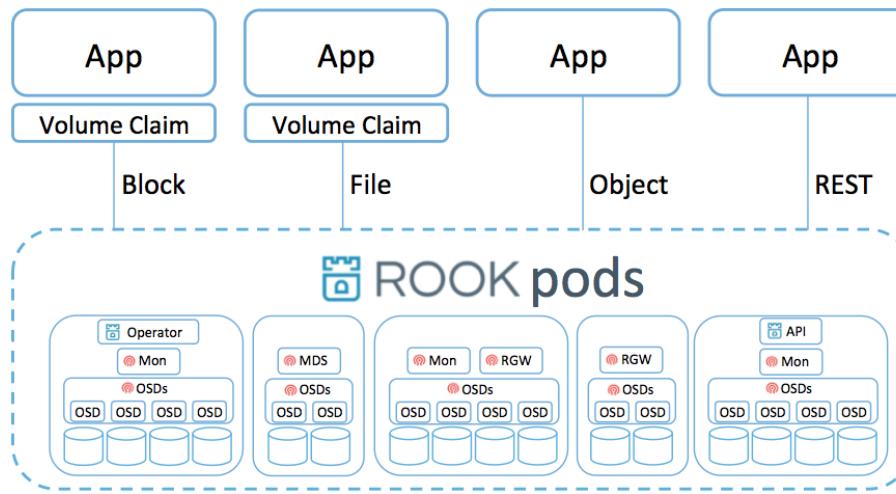


Ceph is a unified, distributed storage system designed for excellent performance, reliability and scalability.

Ceph provide Ceph Object Storage and/or Ceph Block Device services to Cloud Platforms, deploy a Ceph Filesystem or use Ceph for another purpose, all Ceph Storage Cluster deployments begin with setting up each Ceph Node, your network, and the Ceph Storage Cluster.

A Ceph Storage Cluster requires at least one Ceph Monitor, Ceph Manager, and Ceph OSD (Object Storage Daemon). The Ceph Metadata Server is also required when running Ceph Filesystem clients.

Storage: Ceph and Rook (2/2)



Rook is an open source cloud-native **Ceph storage orchestrator** for Kubernetes, providing the platform, framework, and support for a diverse set of storage solutions to natively integrate with cloud-native environments.

Rook turns storage software into self-managing, self-scaling, and self-healing storage services. It does this by automating deployment, bootstrapping, configuration, provisioning, scaling, upgrading, migration, disaster recovery, monitoring, and resource management. Rook uses the facilities provided by the underlying cloud-native container management, scheduling and orchestration platform to perform its duties.

WordPress: Sample Web Application



WORDPRESS

A screenshot of the WordPress dashboard titled "Restaurant WordPress". The dashboard features a sidebar with links like Home, Appearance, Posts, Media, Pages, Comments, Widgets, and Settings. The main area displays "Right Now" stats: 2 Posts, 1 Page, 3 Categories, 6 Comments, 0 Pending, 0 Approved, and 52 Tags. It also shows "Recent Activity" with 2 Widgets and a message from "John Doe" in "Addis Chat" about a dish. The sidebar includes sections for "Recent Posts" (with a link to "Dashboard"), "Recent Comments" (with a link to "Comments"), and "Recent Drafts" (with a link to "Drafts").

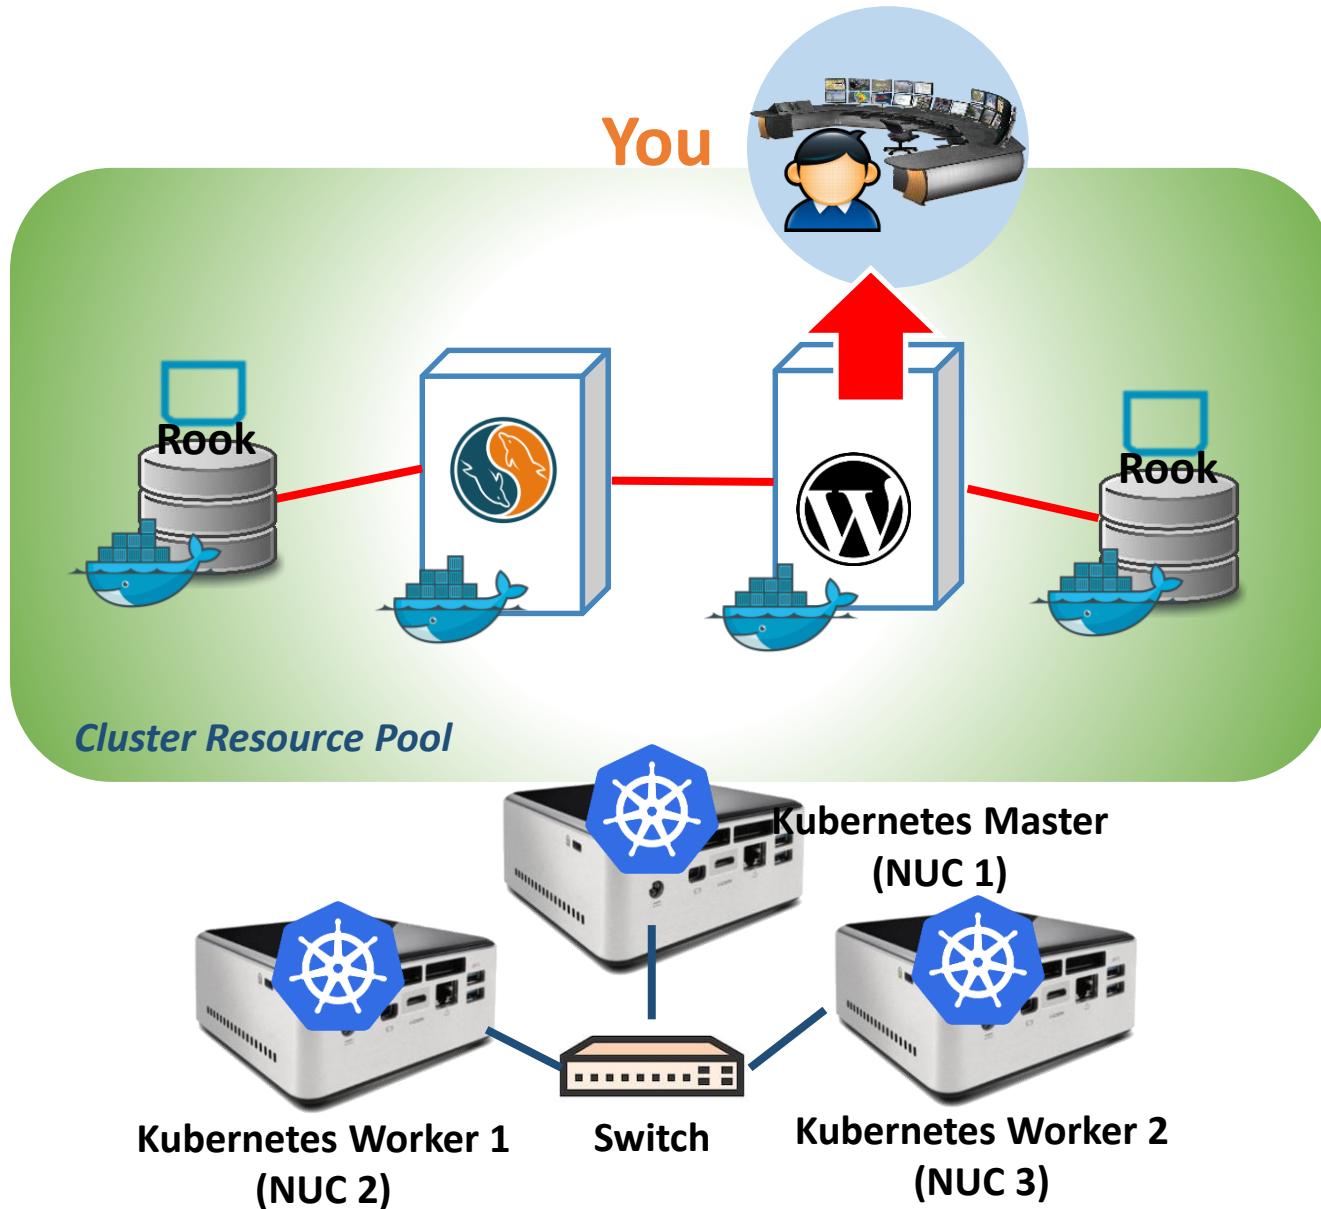
WordPress is the most popular online Open source publishing platform, currently powering more than 28% of the web. Used by more than 60 million websites,[5] including 30.6% of the top 10 million websites as of April 2018, WordPress is the most popular website management system in use.

You can start a blog or build a website in seconds without any technical knowledge!

Practice



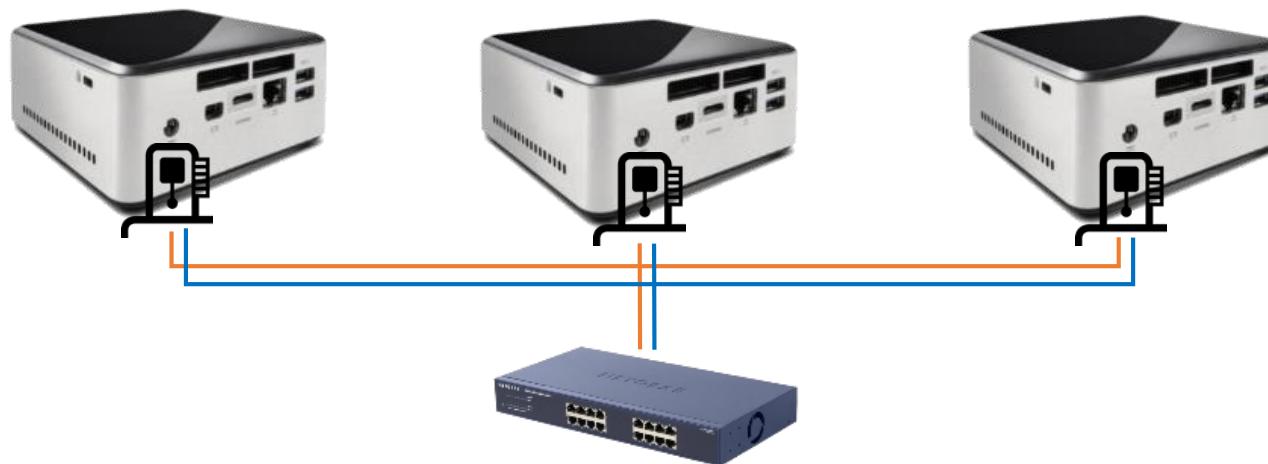
Cluster Lab: Concept



#0 - Lab Preparation (1/4)

Wired connection

NAME: NUC5i5MYHE (NUC PC)
CPU: i5-5300U @2.30GHz
CORE: 4
Memory: 16GB DDR3
HDD: 94GB



NAME: netgear prosafe 16 port gigabit switch(Switch)
Network Ports: 16 auto-sensing 10/100/1000 Mbps Ethernet ports

#0 - Lab Preparation (2/4)

- Boxes preparation: interconnect NUCs

For **NUC1**



In new terminal

```
$ ssh <nuc2 name>@<nuc2 IP address>
> <nuc2 name>@<nuc2 IP address>'s password : <nuc2 pw>
```

In another new terminal

```
$ ssh <nuc3 name>@<nuc3 IP address>
> <nuc3 name>@<nuc3 IP address>'s password : <nuc3 pw>
```

The screenshot shows three terminal windows side-by-side. The window on the left is titled "NUC1 terminal", the middle one "NUC2 terminal", and the right one "NUC3 terminal". The NUC1 terminal shows the command \$ ssh mooc@203. The NUC2 terminal shows the password prompt 'mooc@203. 's password:'. The NUC3 terminal shows the Ubuntu 16.04.2 LTS login screen with the message "Welcome to Ubuntu 16.04.2 LTS (GNU/Linux 4.8.0-36-generic x86_64)". Below the login screen, it lists documentation, management, and support links, followed by package update information: "636 packages can be updated." and "374 updates are security updates.". At the bottom, it shows the last login details: "Last login: Sun Nov 25 14:45:29 2018 from 203.237.53.71".

In NUC1

#0 - Lab Preparation (3/4)

- Boxes preparation: configure hostname



1. From **NUC 1** :

```
$ sudo hostname nuc01
```

2. From **NUC 2** :

```
$ sudo hostname nuc02
```

3. From **NUC 3** :

```
$ sudo hostname nuc03
```

4. Edit /etc/hosts from **all NUCs** :

```
$ sudo vi /etc/hosts
```

5. Append the following context into /etc/hosts :

```
127.0.0.1      localhost
```

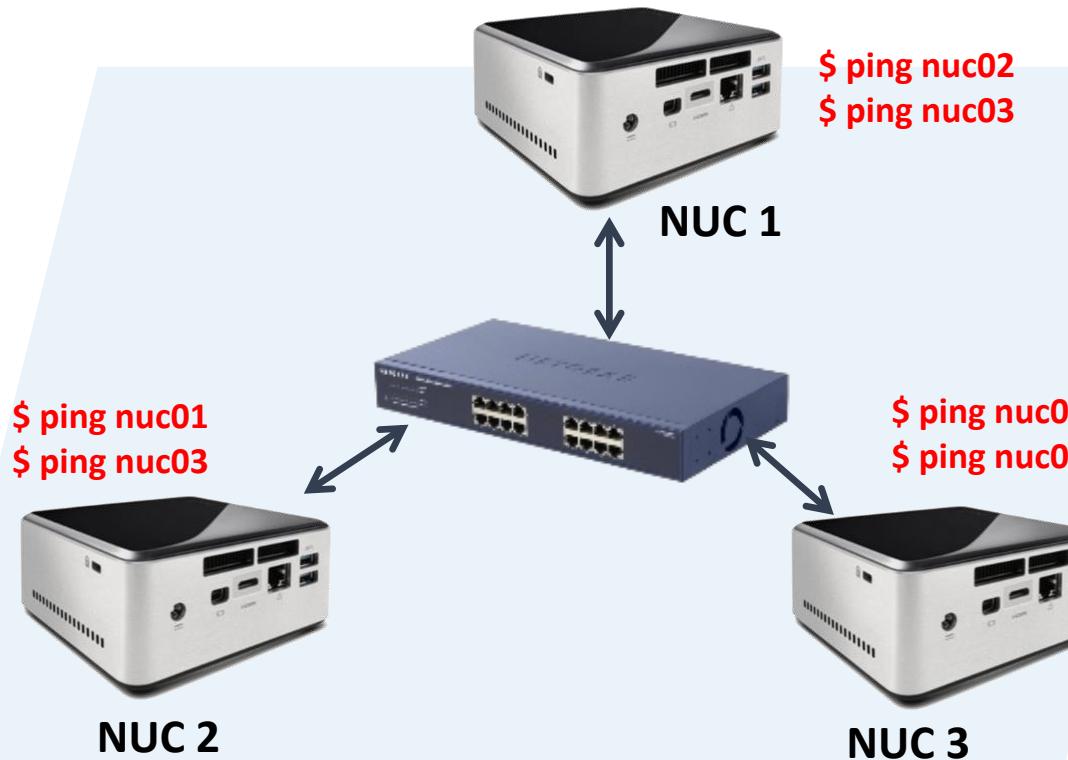
```
IP Address of NUC 1  nuc01
```

```
IP Address of NUC 2  nuc02
```

```
IP Address of NUC 3  nuc03
```

#0 - Lab Preparation (4/4)

- Boxes preparation:
Check full network connectivity between all boxes





#1-1 Preparations for Clustering: Prerequisites Installation

For **All NUCs**

A prerequisite for kubernetes: Docker Installation

```
$ sudo apt-get update  
$ sudo apt-get install apt-transport-https ca-certificates curl software-properties-common  
$ curl -fsSL https://download.docker.com/linux/ubuntu/gpg | sudo apt-key add -  
$ sudo apt-key fingerprint OEBFCD88  
$ sudo add-apt-repository \  
"deb [arch=amd64] https://download.docker.com/linux/ubuntu \  
$(lsb_release -cs) \  
stable"  
$ sudo apt-get update  
$ sudo apt-get install docker-ce
```

A prerequisite for ROOK : xfsprogs installation

```
$ sudo apt-get install xfsprogs
```

#1-2 Preparations for Clustering: Kubernetes Installation



Kubernetes Worker 1
(NUC02)



Kubernetes Master
(NUC01)



Kubernetes Worker 2
(NUC03)

We will deploy kubernetes on All NUCs and configure as above

#1-2 Preparations for Clustering: Kubernetes Installation



For All NUCs

```
$ sudo su  
$ apt-get update && apt-get install -y apt-transport-https curl ipvsadm wget  
$ curl -s https://packages.cloud.google.com/apt/doc/apt-key.gpg | apt-key add -  
$ cat <<EOF | tee /etc/apt/sources.list.d/kubernetes.list  
> deb http://apt.kubernetes.io/ kubernetes-xenial main  
> EOF  
$ curl -s https://packages.cloud.google.com/apt/doc/apt-key.gpg | sudo apt-key add -  
$ apt-get update  
$ apt-get install -y --allow-downgrades kubelet=1.14.1-00 kubeadm=1.14.1-00  
    kubectl=1.14.1-00 kubernetes-cni=0.7.5-00  
$ exit ←
```

Disable root permissions



#1-3 Preparations for Clustering: Kubernetes Configuration

For All NUCs

If swap is not disabled, kubelet service will not start on the masters and nodes. First, you need to edit kubernetes configure file

```
$ sudo vi /etc/systemd/system/kubelet.service.d/10-kubeadm.conf
→ Add [ Environment="KUBELET_EXTRA_ARGS=--fail-swap-on=false" ]
$ sudo systemctl daemon-reload
$ sudo systemctl restart kubelet
```

```
[Service]
Environment="KUBELET_KUBECONFIG_ARGS=--bootstrap-kubeconfig=/etc/kubernetes/bootstrap-kubelet.conf --kubeconfig=/etc/kubernetes/kubelet.conf"
Environment="KUBELET_SYSTEM_PODS_ARGS=--pod-manifest-path=/etc/kubernetes/manifests --allow-privileged=true"
Environment="KUBELET_NETWORK_ARGS=--network-plugin=cni --cni-conf-dir=/etc/cni/net.d --cni-bin-dir=/opt/cni/bin"
Environment="KUBELET_DNS_ARGS=--cluster-dns=10.96.0.10 --cluster-domain=cluster.local"
Environment="KUBELET_AUTHZ_ARGS=--authorization-mode=Webhook --client-ca-file=/etc/kubernetes/pki/ca.crt"
Environment="KUBELET_CADVISOR_ARGS=--cadvisor-port=0"
Environment="KUBELET_CERTIFICATE_ARGS=--rotate-certificates=true --cert-dir=/var/lib/kubelet/pki"
Environment="KUBELET_EXTRA_ARGS=--fail-swap-on=false" Add the line like this
ExecStart=
ExecStart=/usr/bin/kubelet $KUBELET_KUBECONFIG_ARGS $KUBELET_SYSTEM_PODS_ARGS $KUBELET_NETWORK_ARGS $KUBELET_DNS_ARGS $KUBELET_AUTHZ_ARGS $KUBELET_CADVISOR_ARGS $KUBELET_CERTIFICATE_ARGS
```

For All NUCs

Run the following command to disable swap immediately.

```
$ sudo swapoff -a
$ sudo sed -e '/\swapfile/s/^/#/g' -i /etc/fstab
$ sudo sed -e '/\swap\.img/s/^/#/g' -i /etc/fstab
```



#2-1 Clustering: Running Kubernetes Master

For **NUC1**

```
$ sudo systemctl daemon-reload  
$ sudo systemctl restart kubelet  
$ sudo kubeadm reset -f  
$ sudo rm -rf /etc/cni/net.d  
$ sudo ipvsadm --clear  
$ sudo rm -rf /var/lib/rook  
$ sudo kubeadm init --ignore-preflight-errors=all
```

Then you can join any number of worker nodes by running the following on each as root:

```
kubeadm join --token z08hnt.so0bcg051lpe1dk \  
--discovery-token-ca-cert-hash sha256:0060e9bf66f9b38609c3a6196fa0267570902385b77deb8130df3c901e1c3eca
```

You need to copy the command for joining kubernetes nodes (NUC2, NUC3)

Please copy it!

#2-1 Clustering: Running Kubernetes Master



For **NUC1**

Run these commands to make kubectl work for your non-root user.
When you run these commands, make sure that you are not in root

```
$ mkdir -p $HOME/.kube  
$ sudo cp -i /etc/kubernetes/admin.conf $HOME/.kube/config  
$ sudo chown $(id -u):$(id -g) $HOME/.kube/config  
$ kubectl taint nodes --all node-role.kubernetes.io/master-
```

#2-2 Clustering: Joining k8s nodes to k8s master



You can join NUC2, NUC3 to k8s master(NUC1) by following commands

For **NUC2, NUC3**

```
$ sudo systemctl daemon-reload
$ sudo systemctl restart kubelet
$ sudo kubeadm reset -f
$ sudo rm -r /etc/cni/net.d
$ sudo ipvsadm --clear
$ sudo rm -rf /var/lib/rook
$ sudo kubeadm join NUC1 IP:6443 --token YOUR TOKEN --discovery-token-ca-cert-hash
YOUR HASH --ignore-preflight-errors=all → Paste the command you just copied
```

If ipvsadm is not installed,
\$ sudo apt-get install ipvsadm

```
This node has joined the cluster:
* Certificate signing request was sent to apiserver and a response was received.
* The Kubelet was informed of the new secure connection details.
```

Check all nodes are joined

Run 'kubectl get nodes' on the control-plane to see this node join the cluster.

For **NUC1**

```
$ kubectl get node
```

NAME	STATUS	ROLES	AGE	VERSION
nuc01	NotReady	master	1h	v1.11.2
nuc02	NotReady	<none>	6s	v1.11.2
nuc03	NotReady	<none>	4s	v1.11.2



#2-3 Clustering: k8s Network Plugin Installation

You **MUST** install a pod network add-on so that your pods can communicate with each other. We will use Weave in this lab.

For **NUC1**

Install Weave (pod network add-on)

```
$ kubectl apply -f "https://cloud.weave.works/k8s/net?k8s-version=$(kubectl version | base64 | tr -d '\n')"
```

For **NUC1**

Make sure Weave(network add-on) works

```
$ kubectl get nodes
```

```
$ kubectl get po -n kube-system -o wide
```

NAME	STATUS	ROLES	AGE	VERSION
nuc01	Ready	master	1h	v1.11.2
nuc02	Ready	<none>	5m	v1.11.2
nuc03	Ready	<none>	5m	v1.11.2

NAME	READY	STATUS	RESTARTS	AGE
coredns-78fcdf6894-6zktx	1/1	Running	0	1h
coredns-78fcdf6894-glzvq	1/1	Running	0	1h
etcd-nuc01	1/1	Running	0	1h
kube-apiserver-nuc01	1/1	Running	0	1h
kube-controller-manager-nuc01	1/1	Running	0	1h
kube-proxy-6d846	1/1	Running	0	8m
kube-proxy-6n8qr	1/1	Running	0	1h
kube-proxy-r82r7	1/1	Running	0	8m
kube-scheduler-nuc01	1/1	Running	0	1h
weave-net-bw4dz	2/2	Running	0	3m
weave-net-ns426	2/2	Running	0	3m
weave-net-zbs26	2/2	Running	0	3m



#2-3 Clustering: ROOK Installation

For **NUC1**

Set Privileges on cluster for ROOK

```
$ kubectl create clusterrolebinding permissive-binding \
--clusterrole=cluster-admin \
--user=admin \
--user=kubelet \
--group=system:serviceaccounts
```

For **NUC1**

Install ROOK Storage on your cluster

```
$ cd $HOME
$ git clone --single-branch --branch release-1.2 https://github.com/rook/rook.git
$ cd $HOME/rook/cluster/examples/kubernetes/ceph
$ kubectl create -f common.yaml
$ kubectl create -f operator.yaml
$ kubectl create -f cluster-test.yaml
```



#2-3 Clustering: ROOK Installation

For **NUC1**

Check rook-ceph-pod

```
$ watch kubectl get pod -n rook-ceph
```

For **NUC1**

Install ToolBox and Execute ToolBox

```
$ cd $HOME/rook/cluster/examples/kubernetes/ceph
$ kubectl create -f toolbox.yaml
$ kubectl -n rook-ceph rollout status deploy/rook-ceph-tools
$ kubectl -n rook-ceph exec -it $(kubectl -n rook-ceph get pod -l "app=rook-ceph-tools"
-o jsonpath='{.items[0].metadata.name}') bash
```

<in the toolbox>

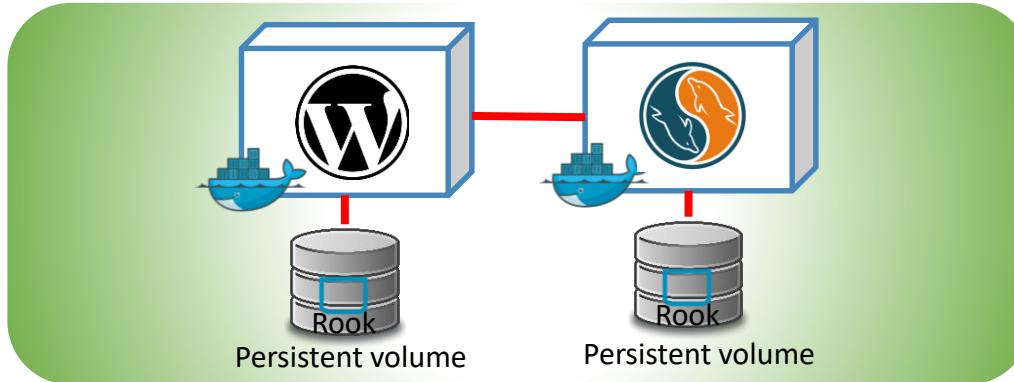
```
$ watch ceph status
$ exit
```

For **NUC1 Add StorageClass**

```
$ kubectl apply -f csi/rbd/storageclass-test.yaml
```

Every 2.0s: kubectl get pod -n rook-ceph						nuc1-NUC1
NAME	READY	STATUS	RESTARTS	AGE		
csi-cephfsplugin-2c2bv	3/3	Running	0	2m21s		
csi-cephfsplugin-5q7z7	3/3	Running	0	2m21s		
csi-cephfsplugin-provisioner-775bb57749-n4h62	4/4	Running	0	2m21s		
csi-cephfsplugin-provisioner-775bb57749-tnbkh	4/4	Running	0	2m21s		
csi-rbdplugin-2qlnw	3/3	Running	0	2m21s		
csi-rbdplugin-m8sgk	3/3	Running	0	2m21s		
csi-rbdplugin-provisioner-65bdfb4644-5dt2p	5/5	Running	0	2m21s		
csi-rbdplugin-provisioner-65bdfb4644-rn4rc	5/5	Running	0	2m21s		
rook-ceph-mon-a-canary-776bd8bd9b-bszzv	1/1	Running	0	83s		
rook-ceph-mon-b-canary-8f9d48567-bw2lb	1/1	Running	0	82s		
rook-ceph-mon-c-canary-7d655d4f7d-8qw6s	1/1	Running	0	82s		
rook-ceph-operator-8644d9bd85-2rann	1/1	Running	0	2m42s		
rook-discover-8mnmk	1/1	Running	0	2m24s		
rook-discover-xd8r2	1/1	Running	0	2m24s		

#2-4 Cluster Validation: Deploy Sample Web application



For **NUC1**

Deploy a Sample Web Application (WordPress) on your cluster.

```
$ kubectl create -f $HOME/rook/cluster/examples/kubernetes/mysql.yaml  
$ kubectl create -f $HOME/rook/cluster/examples/kubernetes/wordpress.yaml
```

For **NUC1**

Check Wordpress containers are started

```
$ watch kubectl get pod
```

Every 2.0s: kubectl get po					
NAME	READY	STATUS	RESTARTS	AGE	
wordpress-595685cc49-ct9f1	1/1	Running	0	50s	
wordpress-mysql-b78774f44-llgst	1/1	Running	0	56s	



#2-4 Cluster Validation: Deploy Sample Web application

For **NUC1**

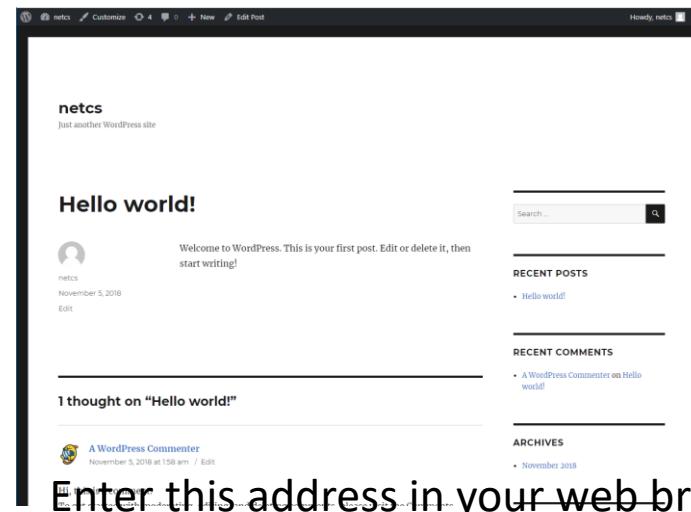
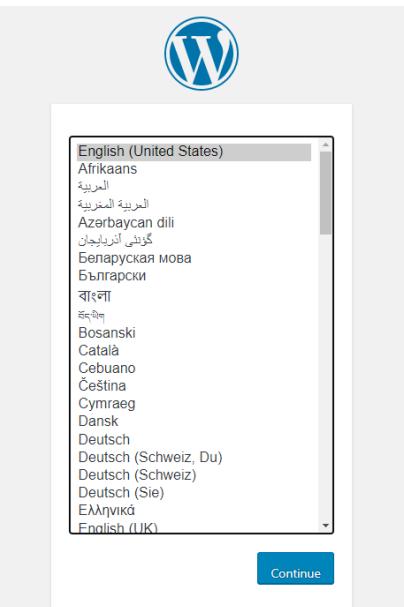
Check exposed port for Wordpress service.

The address of your Wordpress Web is <http://your NUC1 IP: exposed port>

\$ kubectl get svc

exposed port

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
kubernetes	ClusterIP	10.96.0.1	<none>	443/TCP	160m
wordpress	LoadBalancer	10.107.116.108	<pending>	80:31671/TCP	20m
wordpress-mysql	ClusterIP	None	<none>	3306/TCP	20m



Enter this address in your web browser:
<http://your NUC1 IP:exposed port>

Review

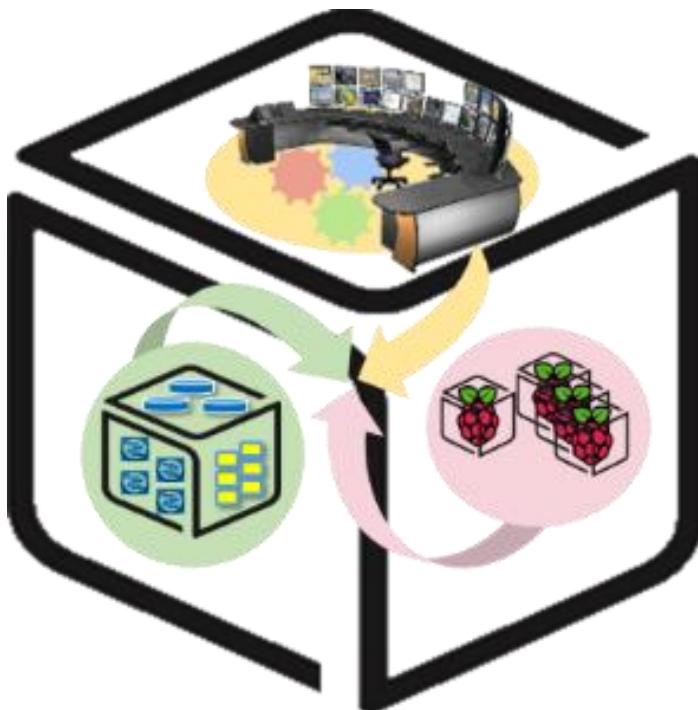


Lab Summary

With Cluster Lab, you have experimented

1. How to physically connect multiple NUCs with a switch to create the targeted physical cluster?
2. How to install (deploy) and configure container-orchestrated software functions to enable the desired cluster role (e.g., HPC, HPDA, ML/DL, ...)
3. Understanding how container orchestration is working: Can you distinguish Kubernetes master from Kubernetes nodes?

Thank You for Your Attention Any Questions?



mini@smartx.kr